

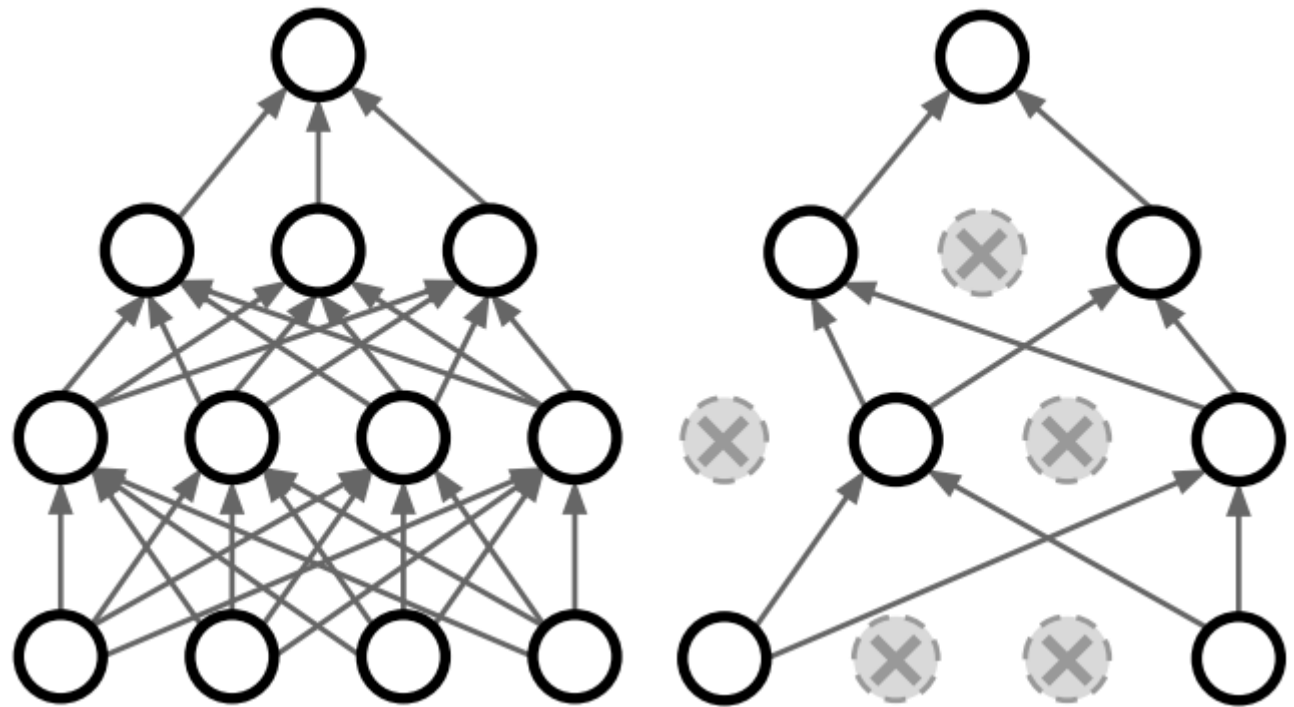
رسالة محمد

Deep Learning

Mohammad Reza Mohammadi
2021

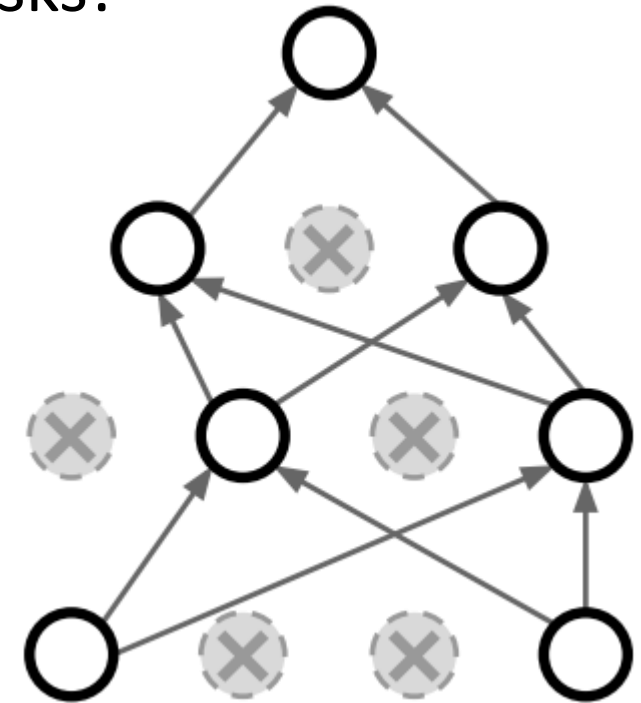
Dropout

- In each forward pass, randomly set some neurons to zero
- Probability of dropping is a hyperparameter; 0.5 is common



Dropout

- Dropout is training a large ensemble of models (that share parameters)
- Each binary mask is one model
- An FC layer with 4096 units has $2^{4096} \sim 10^{1233}$ possible masks!

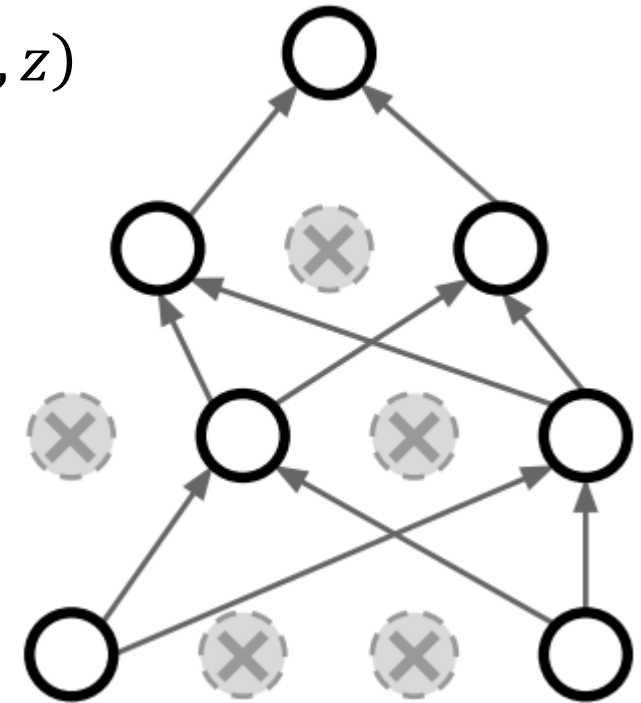


Dropout: Test time

- Dropout makes our output random!
- Want to “average out” the randomness at test-time

$$y = f(x) = E_z[f(x, z)] = \sum_z p(z) f(x, z)$$

- But this summation is very time-consuming



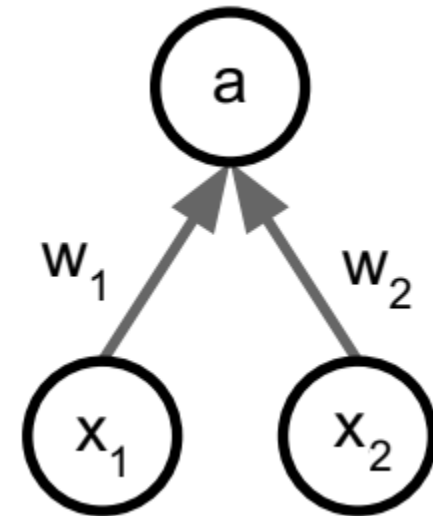
Dropout: Test time

- Want to approximate the expected value
- Consider a single neuron

$$y = f(x) = E_z[f(x, z)] = \sum_z p(z) f(x, z)$$

$$\begin{aligned} E[a] &= p^2(w_1x_1 + w_2x_2) + p(1-p)(w_1x_1 + w_20) \\ &\quad + (1-p)p(w_10 + w_2x_2) + (1-p)^2(w_10 + w_20) \\ &= p(w_1x_1 + w_2x_2) \end{aligned}$$

- At test time, multiply by dropout probability



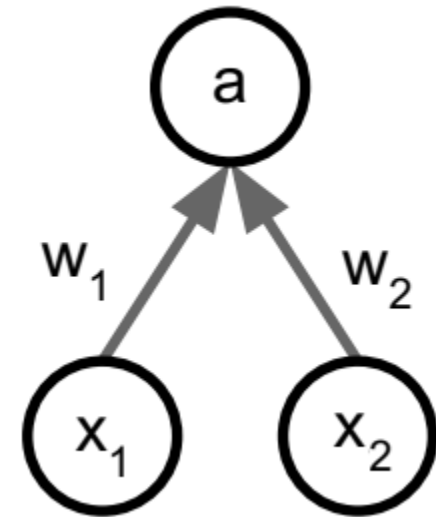
Dropout: Test time

- Want to approximate the expected value
- Consider a single neuron

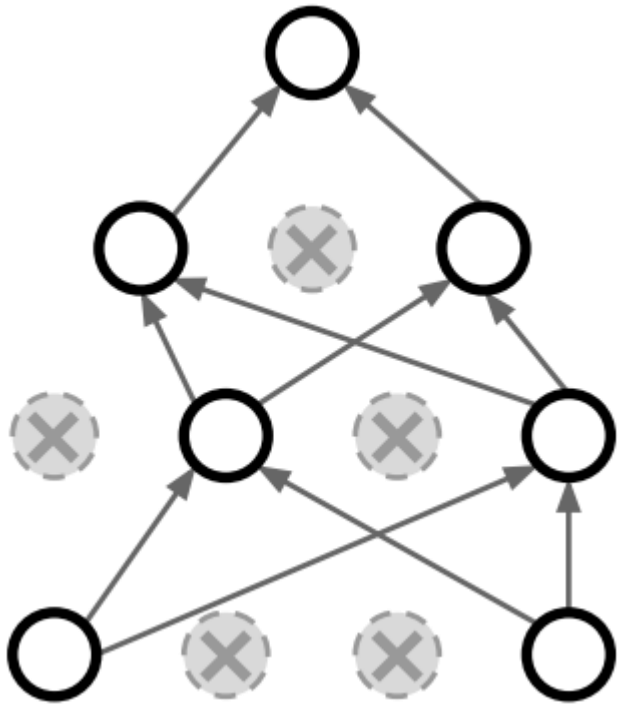
$$y = f(x) = E_z[f(x, z)] = \sum_z p(z) f(x, z)$$

$$\begin{aligned} E[a] &= p^2(w_1x_1 + w_2x_2) + p(1-p)(w_1x_1 + w_20) \\ &\quad + (1-p)p(w_10 + w_2x_2) + (1-p)^2(w_10 + w_20) \\ &= p(w_1x_1 + w_2x_2) \end{aligned}$$

```
def predict(X):  
    # ensembled forward pass  
    H1 = np.maximum(0, np.dot(W1, X) + b1) * p # NOTE: scale the activations  
    H2 = np.maximum(0, np.dot(W2, H1) + b2) * p # NOTE: scale the activations  
    out = np.dot(W3, H2) + b3
```



Dropout



```
p = 0.5 # probability of keeping a unit active. higher = less dropout
def train_step(X):
```

```
    # forward pass for example 3-layer neural network
```

```
    H1 = np.maximum(0, np.dot(W1, X) + b1)
```

```
    U1 = (np.random.rand(*H1.shape) < p) # first dropout mask.
```

```
    H1 *= U1 # drop!
```

```
    H2 = np.maximum(0, np.dot(W2, H1) + b2)
```

```
    U2 = (np.random.rand(*H2.shape) < p) # second dropout mask.
```

```
    H2 *= U2 # drop!
```

```
    out = np.dot(W3, H2) + b3
```

drop in forward pas

```
    # backward pass: compute gradients... (not shown)
```

```
    # perform parameter update... (not shown)
```

```
def predict(X):
```

```
    # ensembled forward pass
```

```
    H1 = np.maximum(0, np.dot(W1, X) + b1)
```

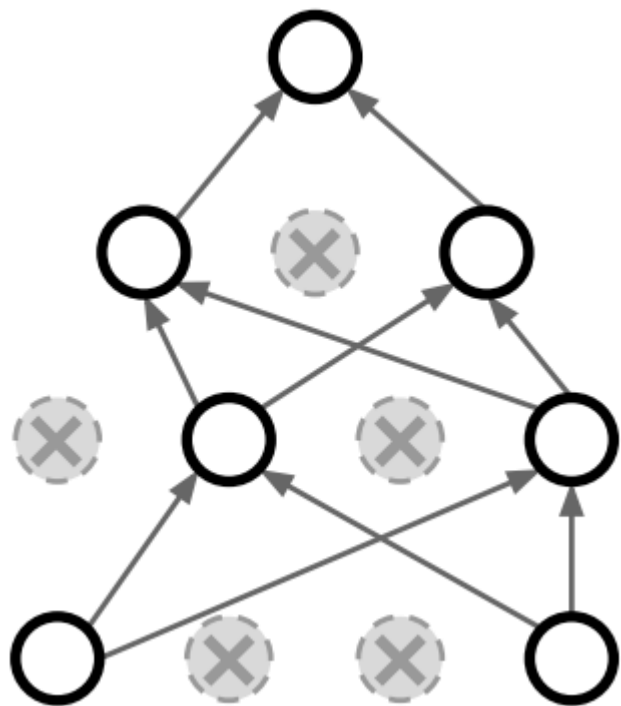
```
    H2 = np.maximum(0, np.dot(W2, H1) + b2)
```

```
    out = np.dot(W3, H2) + b3
```

scale at test time

```
* p # scale the activations
* p # scale the activations
```


Inverted Dropout



```
p = 0.5 # probability of keeping a unit active. higher = less dropout
```

```
def train_step(X):
```

```
    # forward pass for example 3-layer neural network
```

```
    H1 = np.maximum(0, np.dot(W1, X) + b1)
```

```
    U1 = (np.random.rand(*H1.shape) < p) / p # first dropout mask. Notice /p!
```

```
    H1 *= U1 # drop!
```

```
    H2 = np.maximum(0, np.dot(W2, H1) + b2)
```

```
    U2 = (np.random.rand(*H2.shape) < p) / p # second dropout mask. Notice /p!
```

```
    H2 *= U2 # drop!
```

```
    out = np.dot(W3, H2) + b3
```

drop in forward pass

```
    # backward pass: compute gradients... (not shown)
```

```
    # perform parameter update... (not shown)
```

```
def predict(X):
```

```
    # ensembled forward pass
```

```
    H1 = np.maximum(0, np.dot(W1, X) + b1) # no scaling necessary
```

```
    H2 = np.maximum(0, np.dot(W2, H1) + b2)
```

```
    out = np.dot(W3, H2) + b3
```

test time is unchanged

Convolutional Networks

```

early_stopping = keras.callbacks.EarlyStopping(monitor="val_loss",
                                                min_delta=0,
                                                patience=5,
                                                restore_best_weights=True)

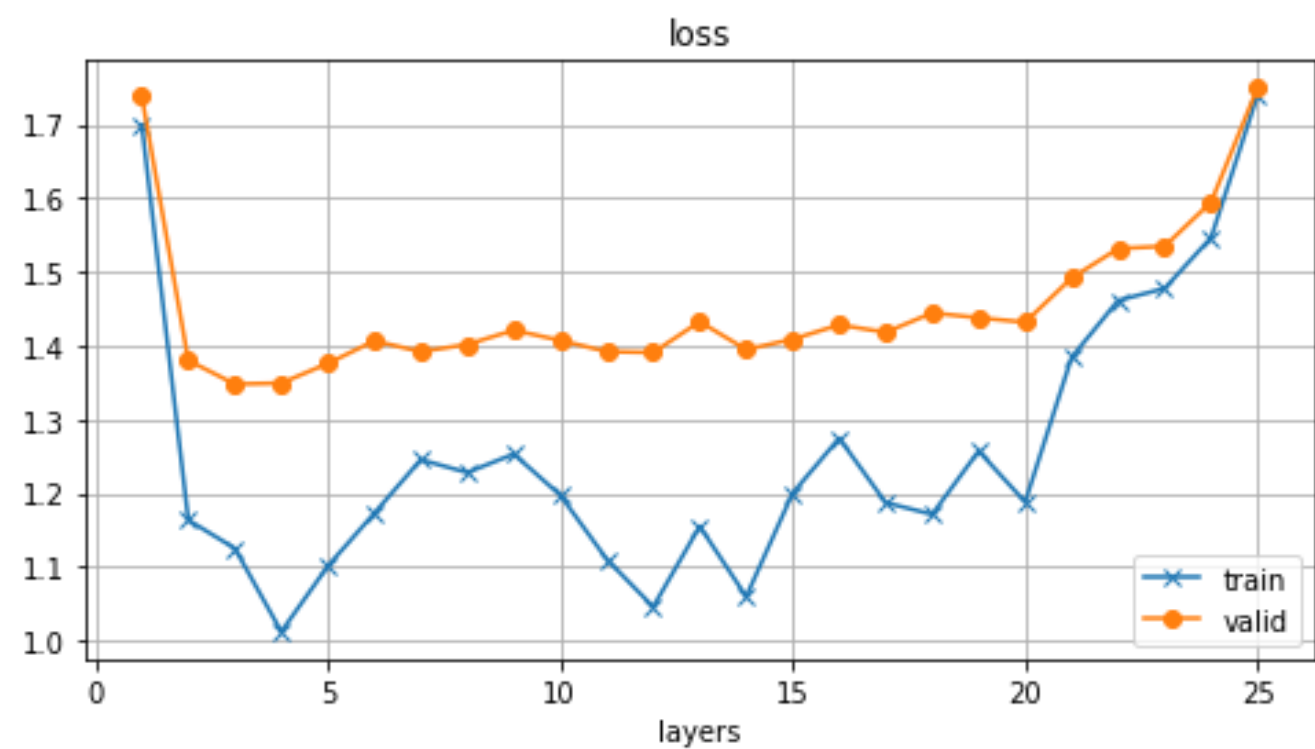
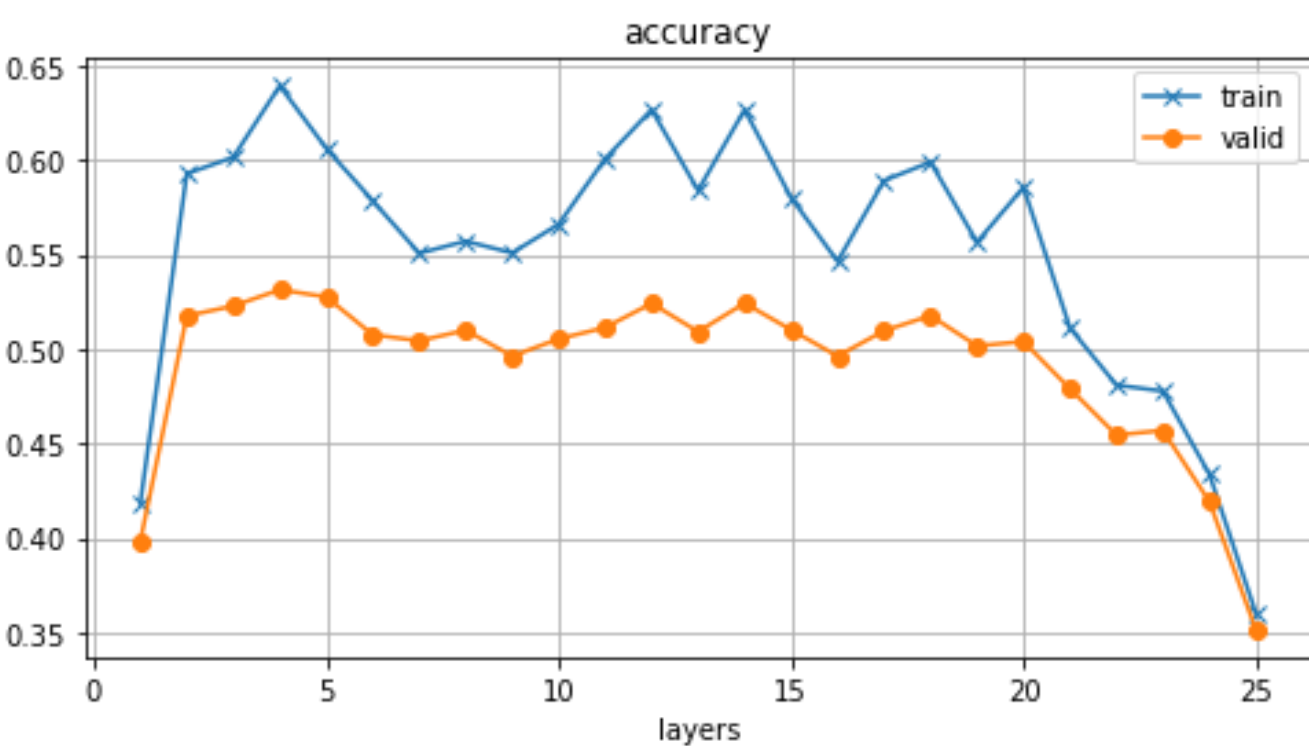
for idx, num_layers in enumerate(range(25)):
    # define model
    model = keras.Sequential()
    model.add(keras.layers.Input(shape=x_train[0].shape))
    model.add(keras.layers.Flatten())
    for l in range(num_layers):
        model.add(keras.layers.Dense(units=512, activation='elu'))
    model.add(keras.layers.Dense(units=num_classes, activation='softmax'))

    # compile model
    model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

    # train model
    history = model.fit(x_train, y_train,
                        batch_size=256,
                        epochs=100,
                        validation_data=(x_test, y_test),
                        callbacks=[early_stopping])

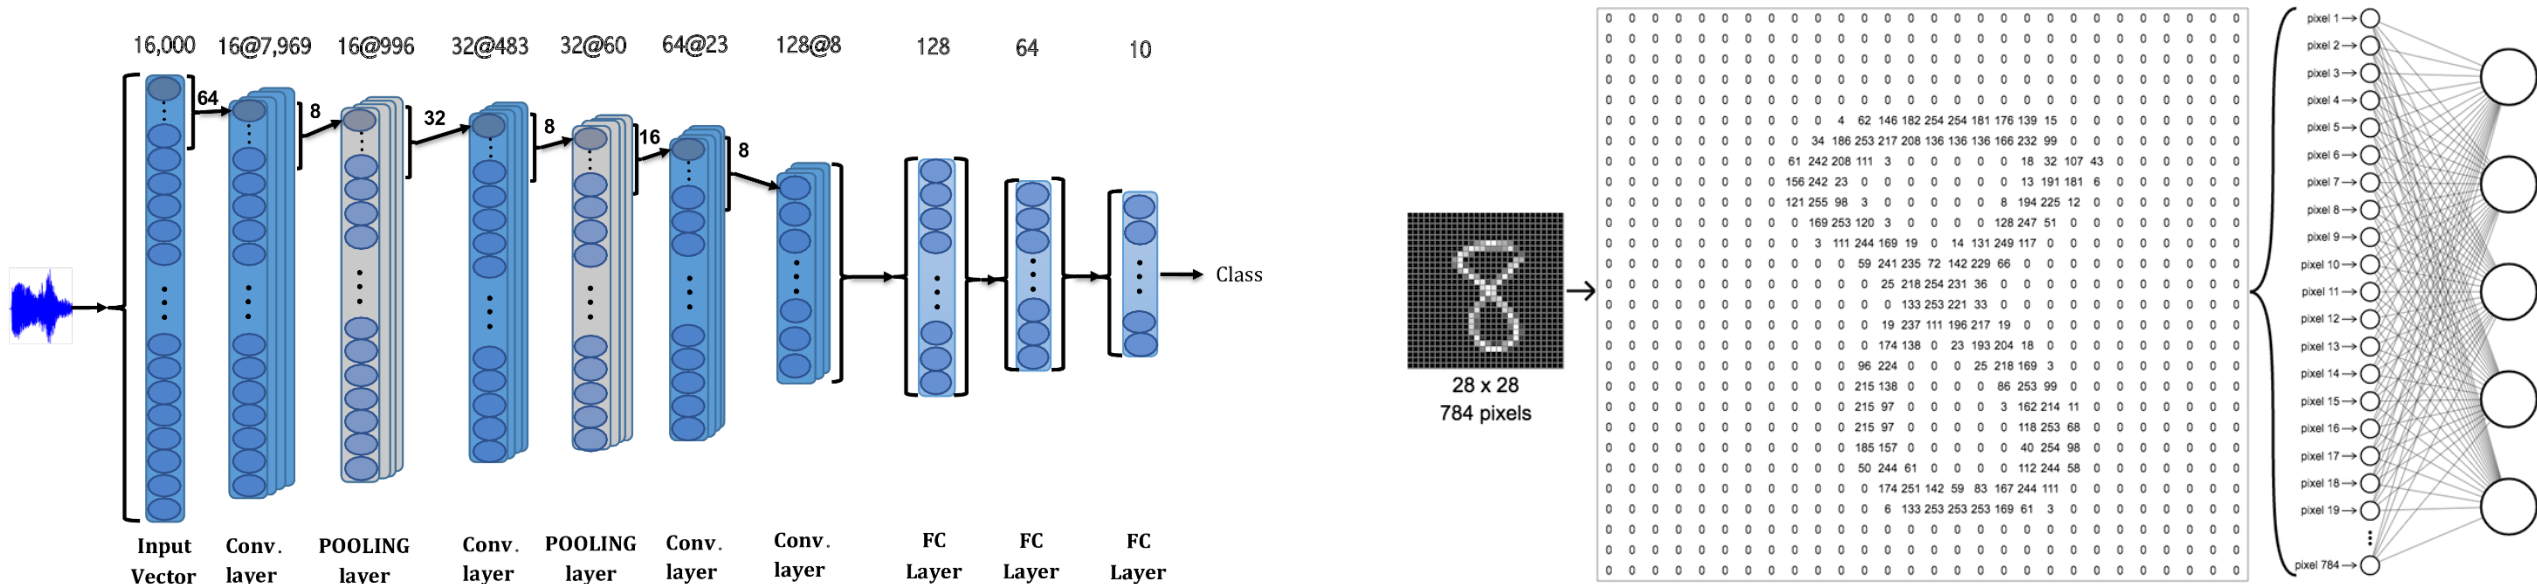
    #
    train_loss[idx], train_acc[idx] = model.evaluate(x_train, y_train, verbose=0)
    valid_loss[idx], valid_acc[idx] = model.evaluate(x_test, y_test, verbose=0)

```



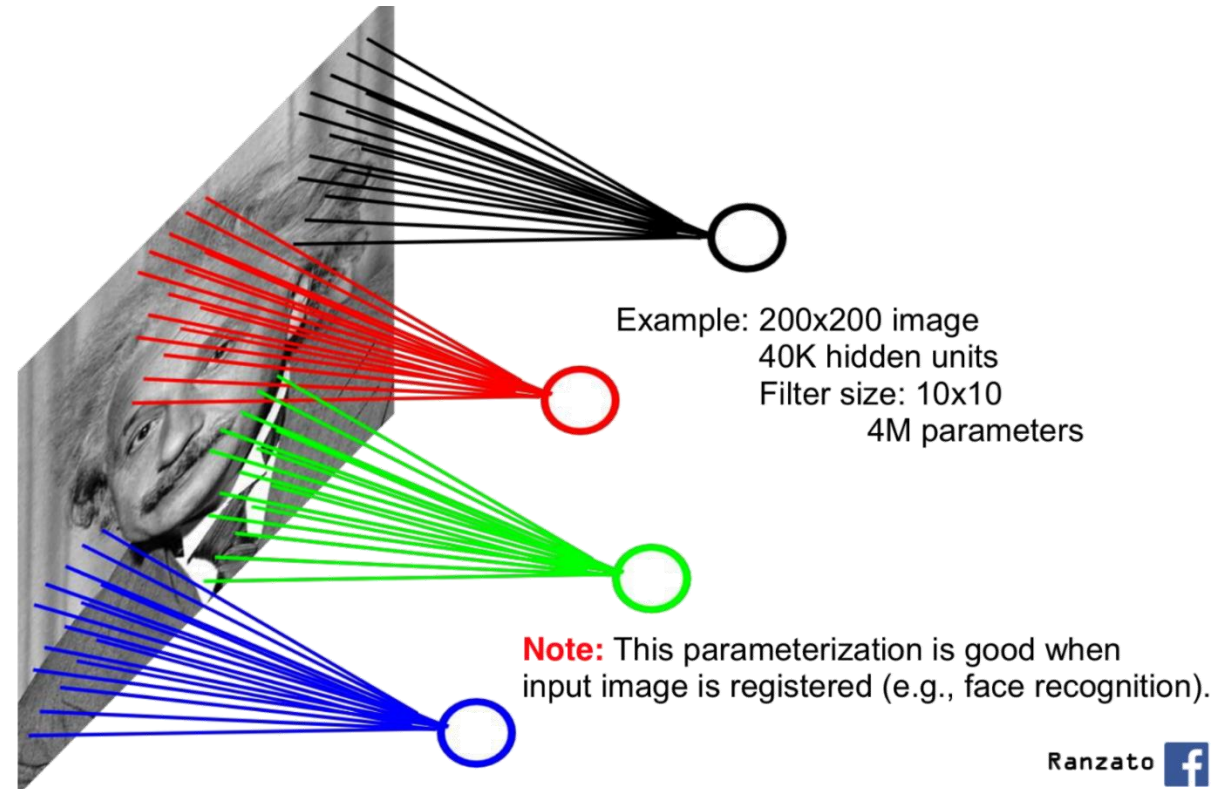
Convolutional Networks

- CNNs, are a specialized kind of NN for processing data that has a known, grid-like topology
- Convolution is a specialized kind of linear operation



Locally Connected Layers

- Many features that the human eye can easily detect are local features
- We can detect edges, textures, and even shapes using pixel intensities in a small region of an image
- If we wanted to detect a feature, we can use the same detector on the bottom-left corner of an image and on the top right of the image
- We can reuse the same weights everywhere in the image
 - **weight sharing**

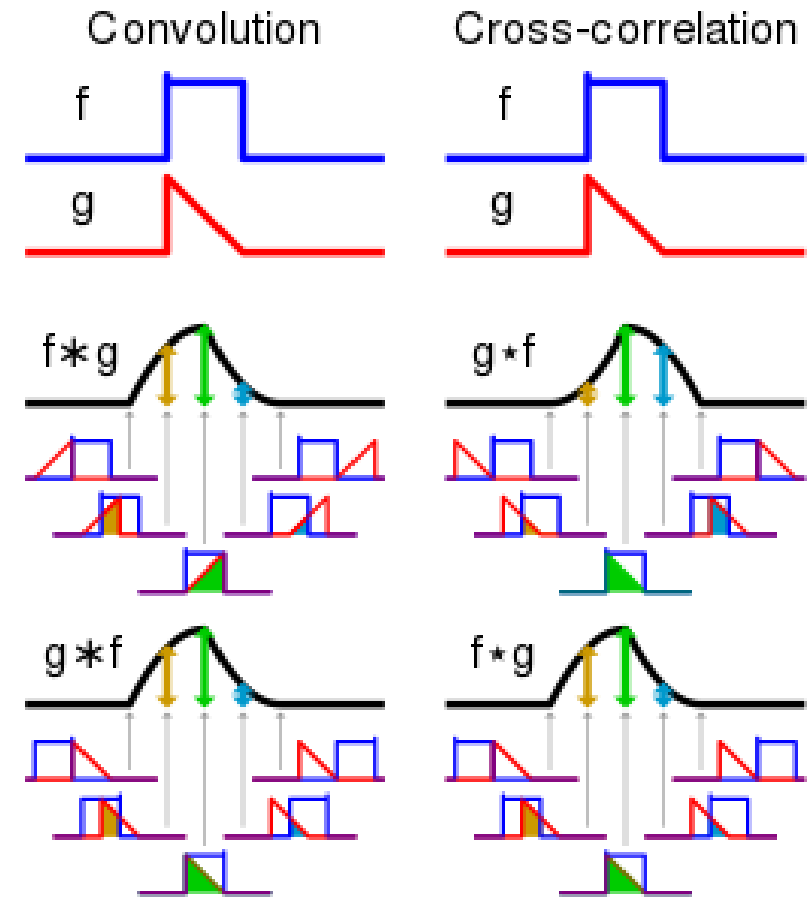


Convolution vs Correlation

- Many ML libraries implement cross-correlation but call it convolution!
- The learning algorithm will learn the appropriate values of the kernel in the appropriate place

$$S(i, j) = (I \star K)(i, j) = \sum_m \sum_n I(i + m, j + n) K(m, n)$$

$$S(i, j) = (I * K)(i, j) = \sum_m \sum_n I(i - m, j - n) K(m, n)$$



Convolution



3_0	3_1	2_2	1	0
0_2	0_2	1_0	3	1
3_0	1_1	2_2	2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

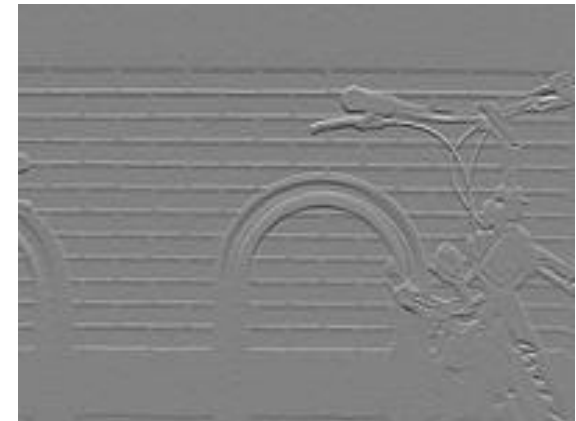
G_y

+1	0	-1
+2	0	-2
+1	0	-1



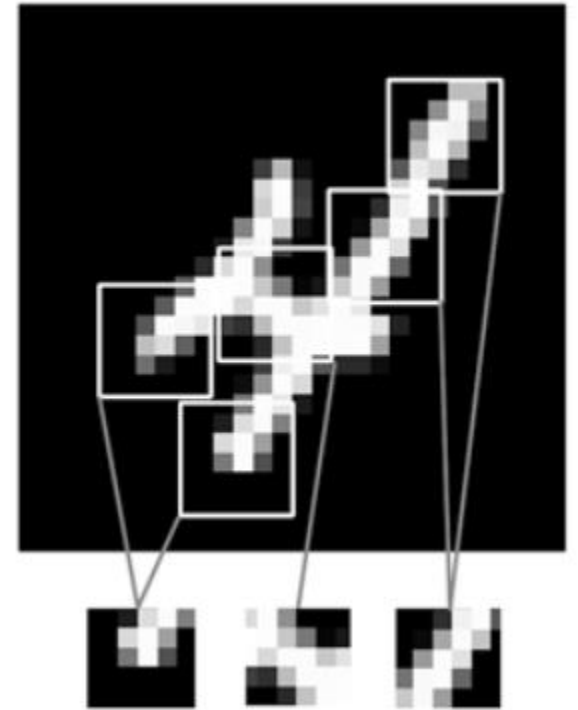
G_x

+1	+2	+1
0	0	0
-1	-2	-1



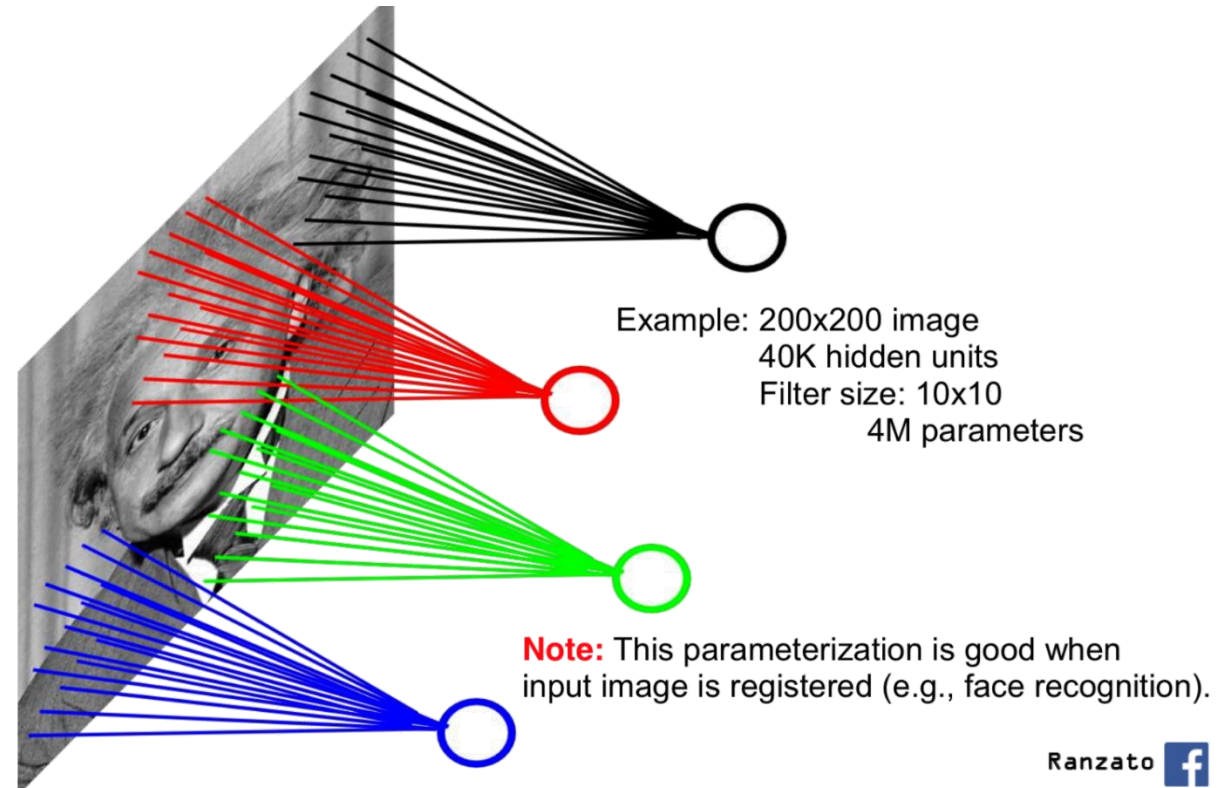
Dense layer vs Convolution layer

- Dense layers learn global patterns in their input feature space (for example, for a MNIST digit, patterns involving all pixels)
- Convolution layers learn local patterns



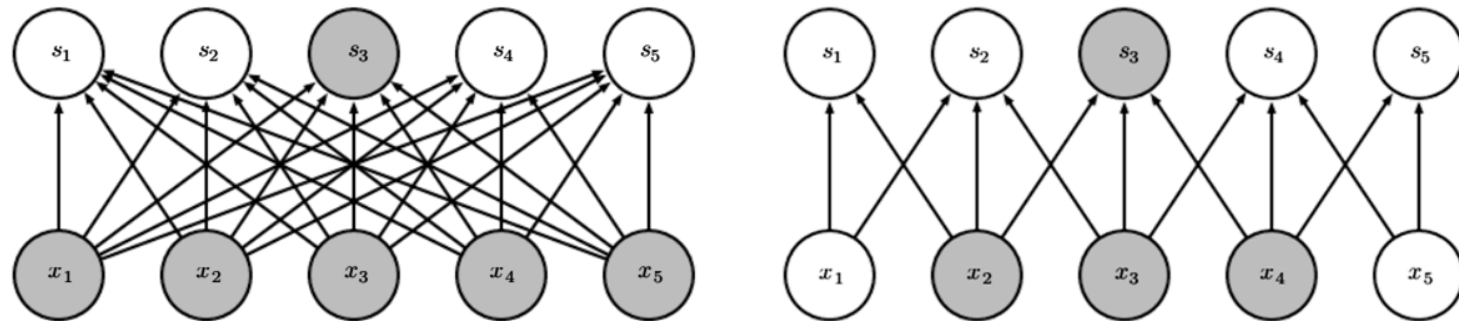
Convolution

- Sparse interactions
- Parameter sharing
- Equivariant representations
- Working with inputs of variable size



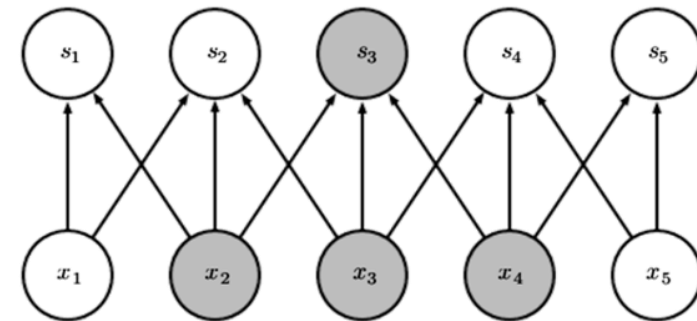
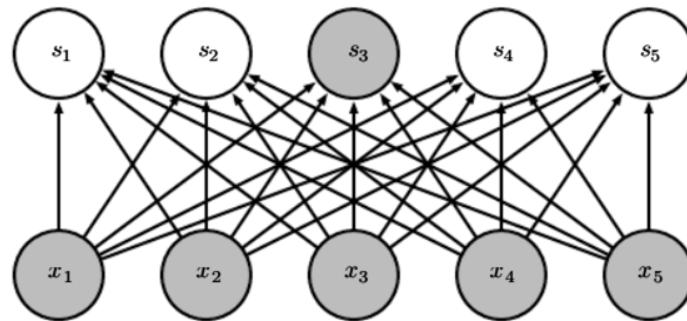
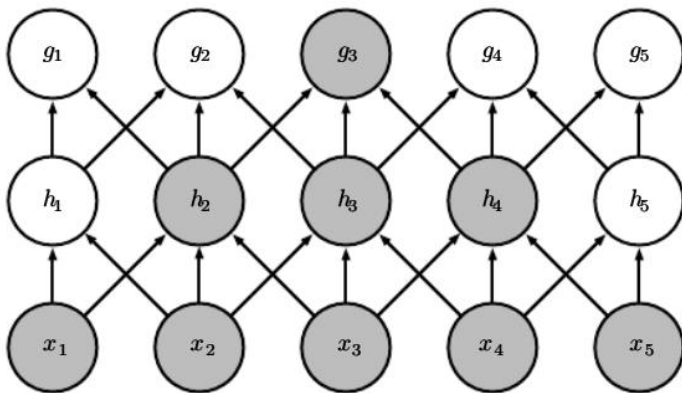
Sparse interactions

- In traditional NN layers, every output unit interacts with every input unit
- Convolutional networks, typically have sparse interactions
- For example, when processing an image, the input image might have thousands or millions of pixels, but we can detect small, meaningful features such as edges with kernels that occupy only tens or hundreds of pixels
 - We need to store fewer parameters



Sparse interactions

- In a deep convolutional network, units in the deeper layers may indirectly interact with a larger portion of the input
- This allows the network to efficiently describe complicated interactions between many variables by constructing such sparse interactions
- The receptive field of the units in the deeper layers can be very large



Parameter sharing

- Parameter sharing refers to using the same parameter for more than one function in a model
- In a convolutional neural net, each member of the kernel is used at every position of the input
- The parameter sharing means that rather than learning a separate set of parameters for every location, we learn only one set

