

رَبِّ الْعَالَمِينَ

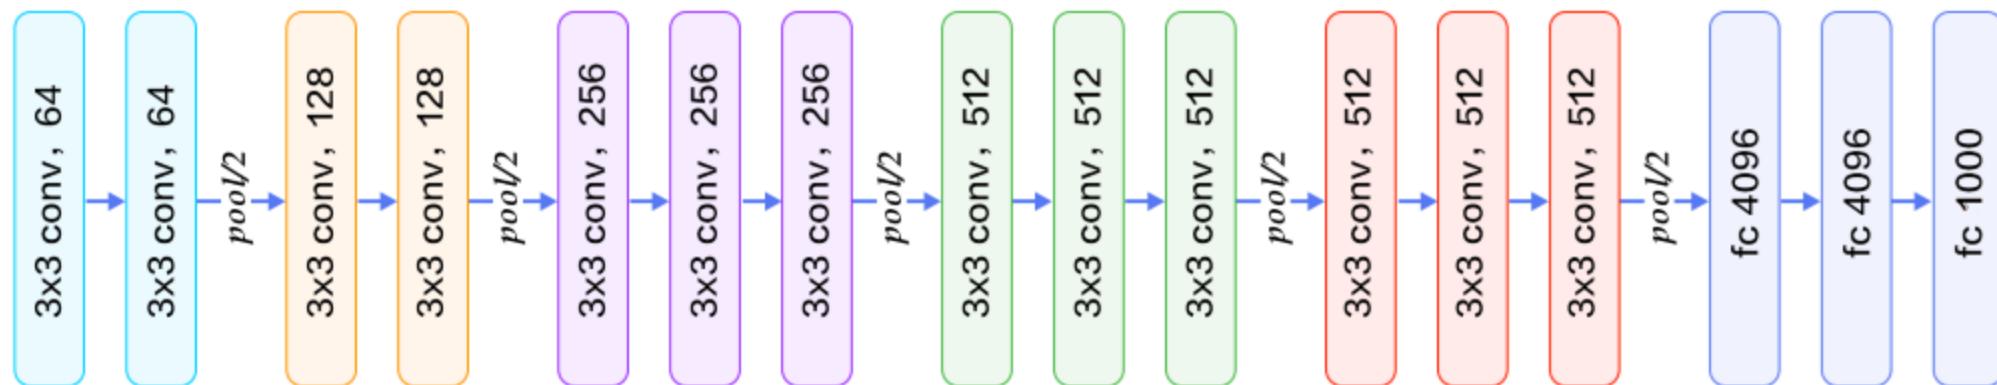
# Deep Learning

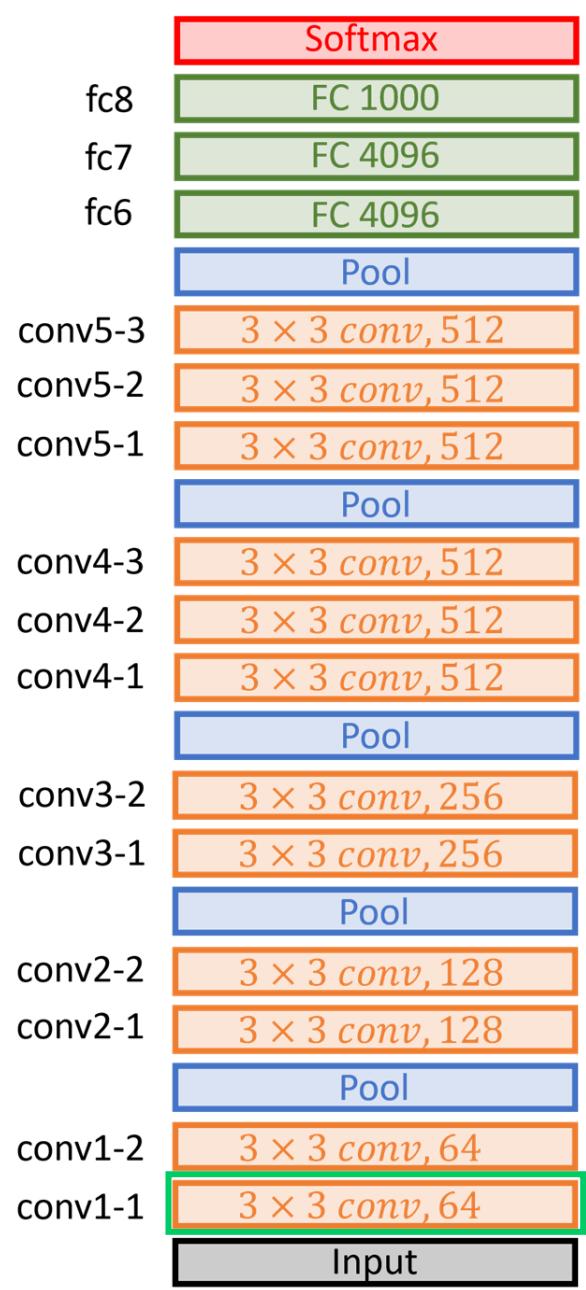
Mohammad Reza Mohammadi

2021

# Visualizing what convnets learn

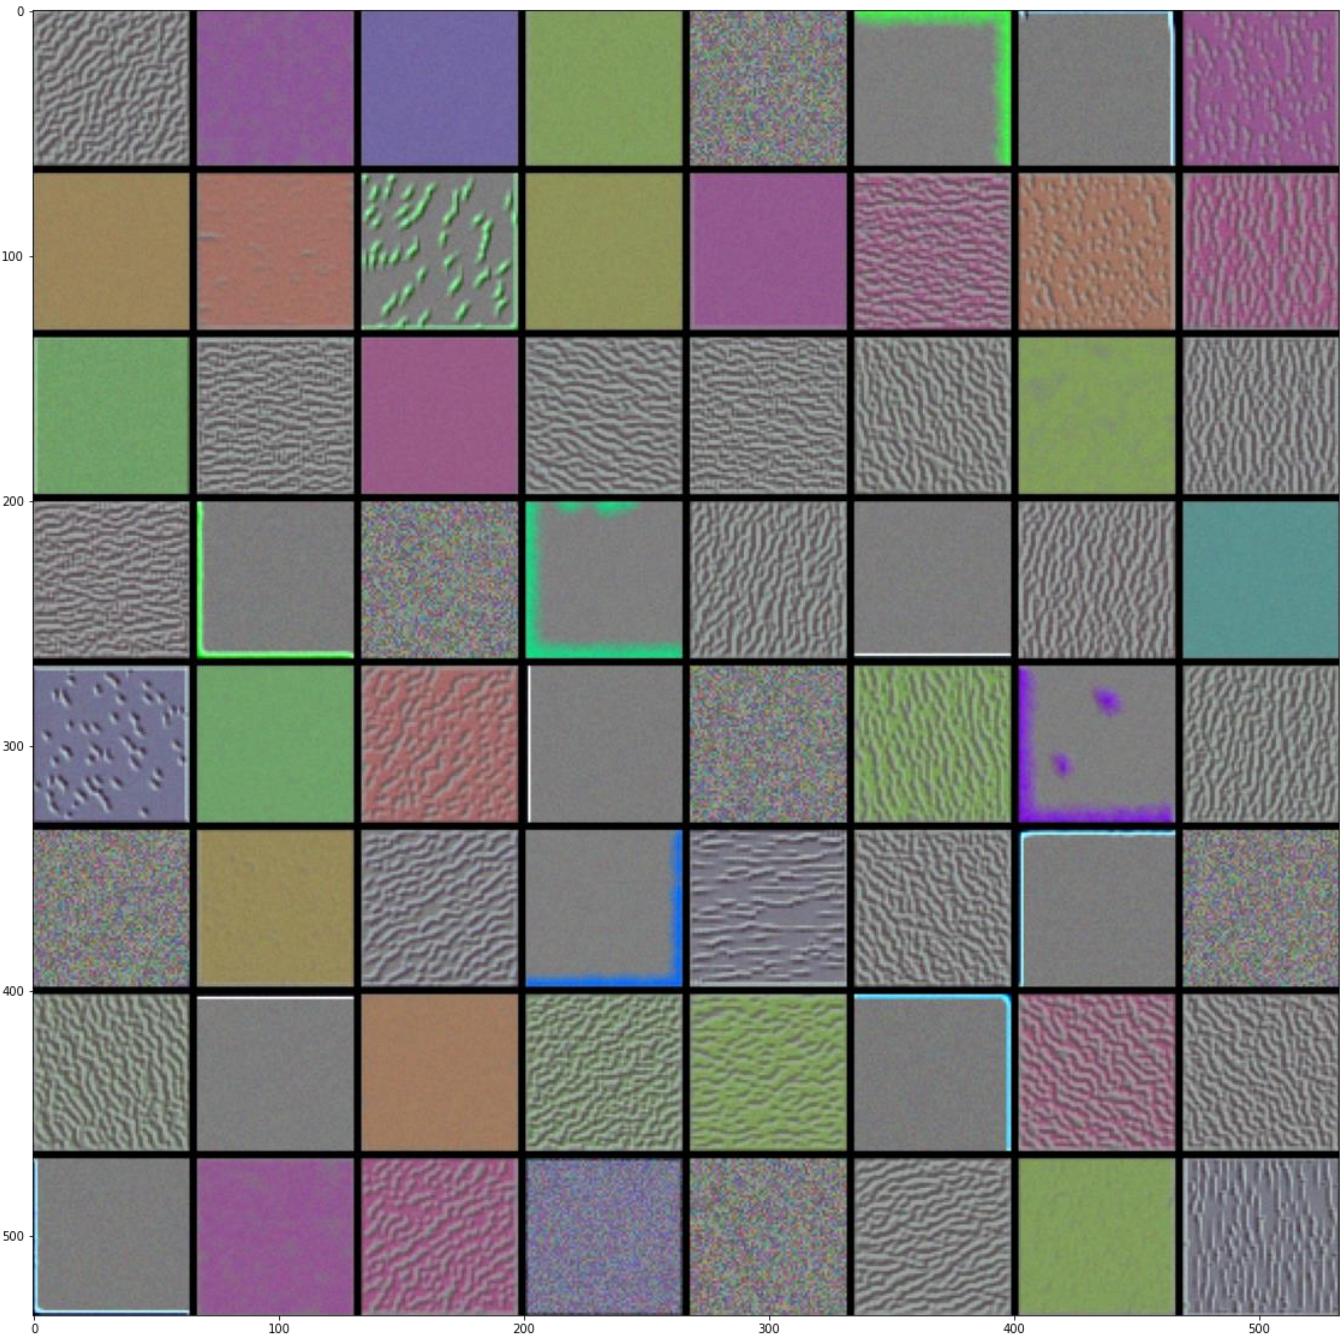
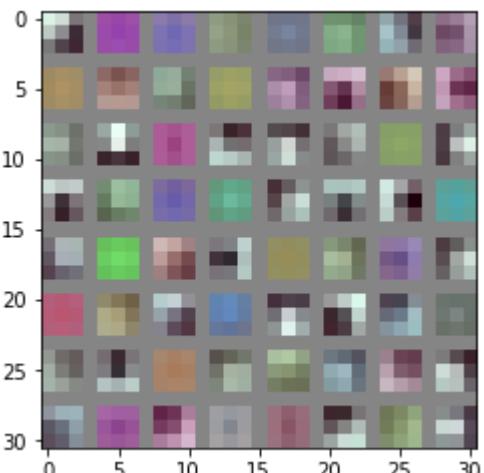
- Different ways to visualize or interpret NN representations, three common ones:
  - Visualizing intermediate convnet outputs (intermediate activations)
  - Visualizing convnets filters
  - Visualizing heatmaps of class activation in an image

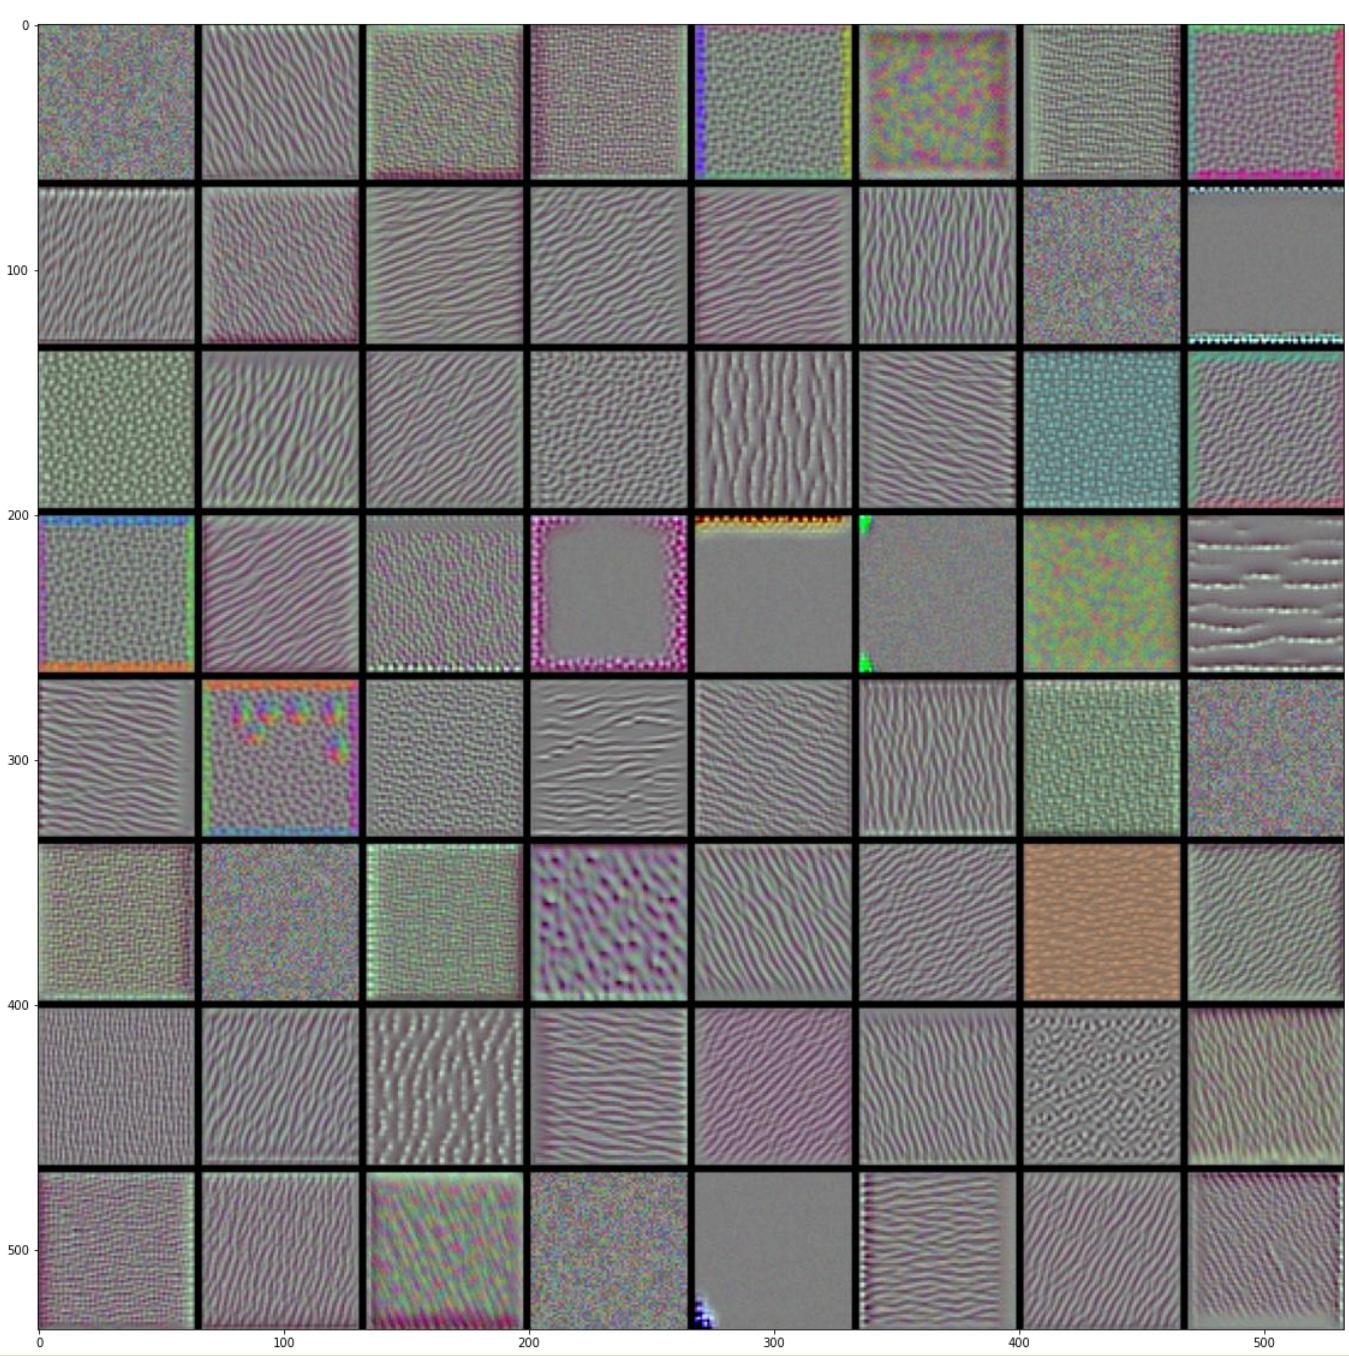
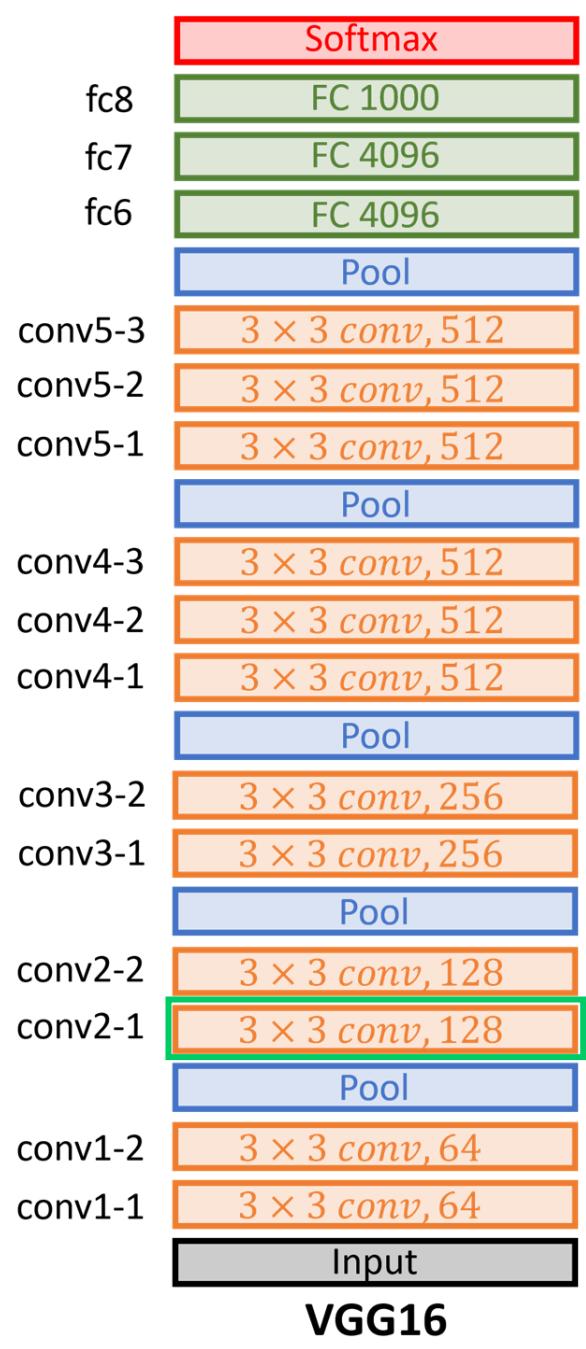


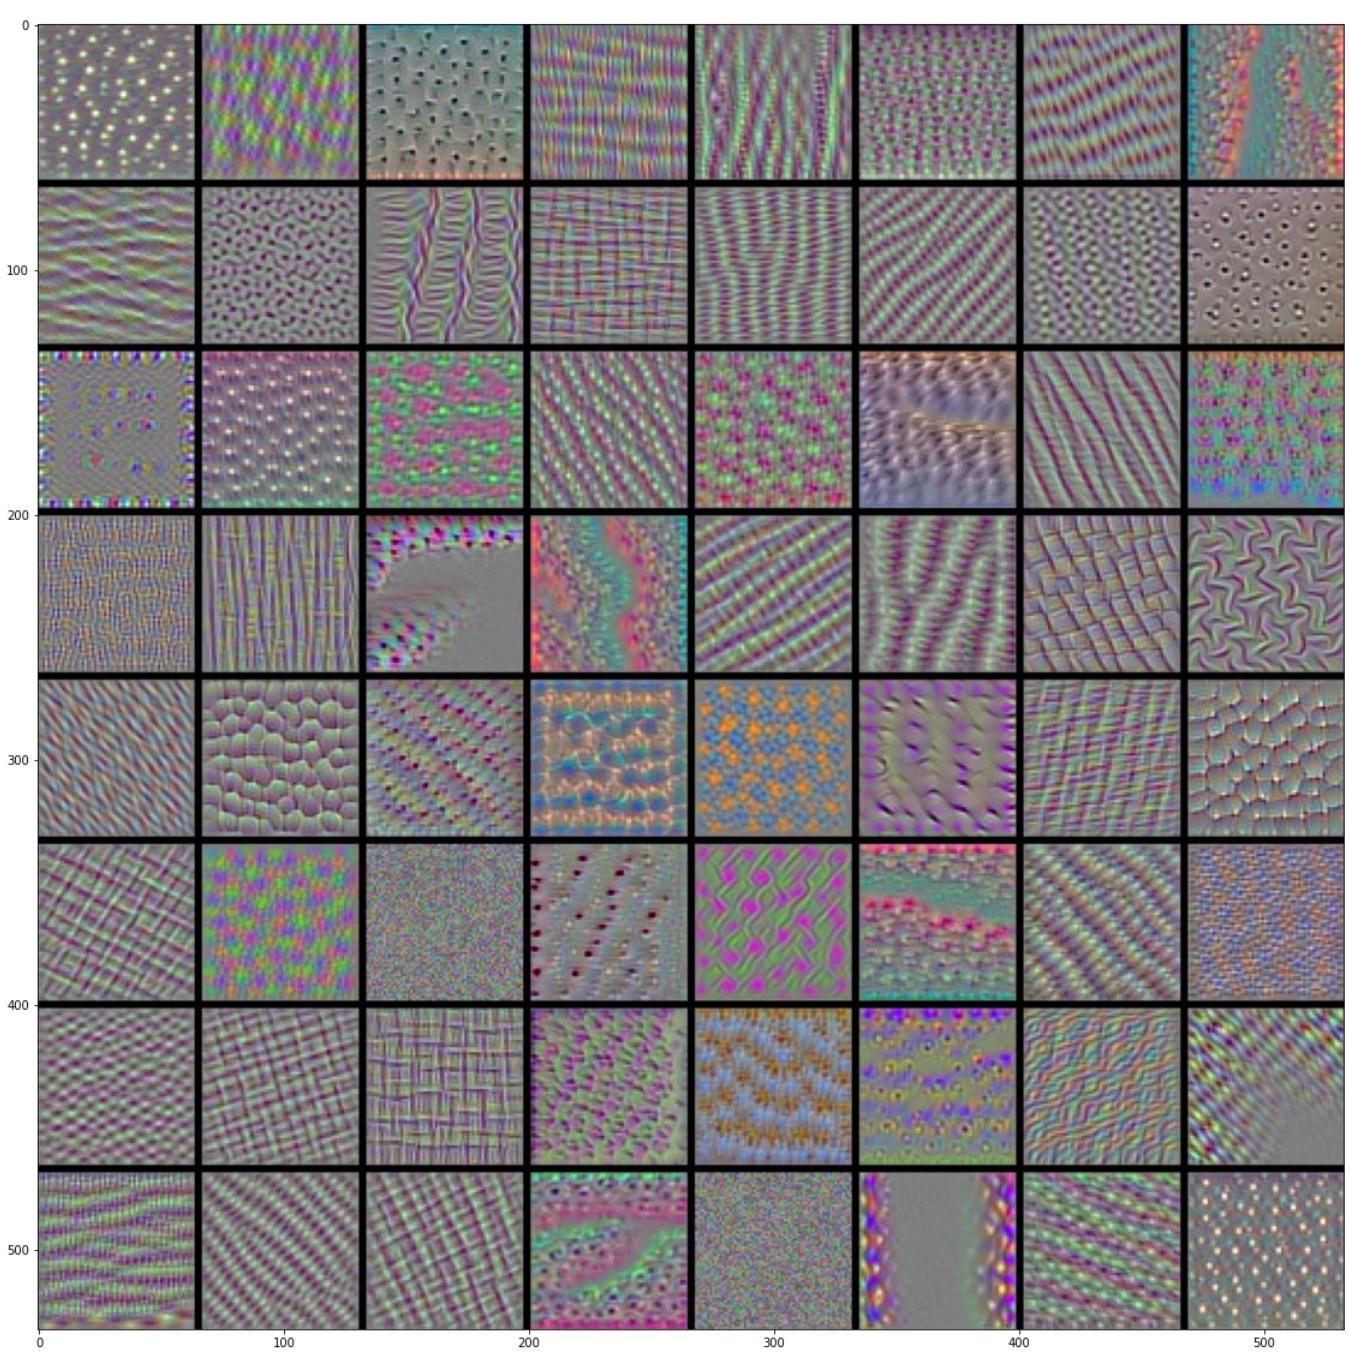
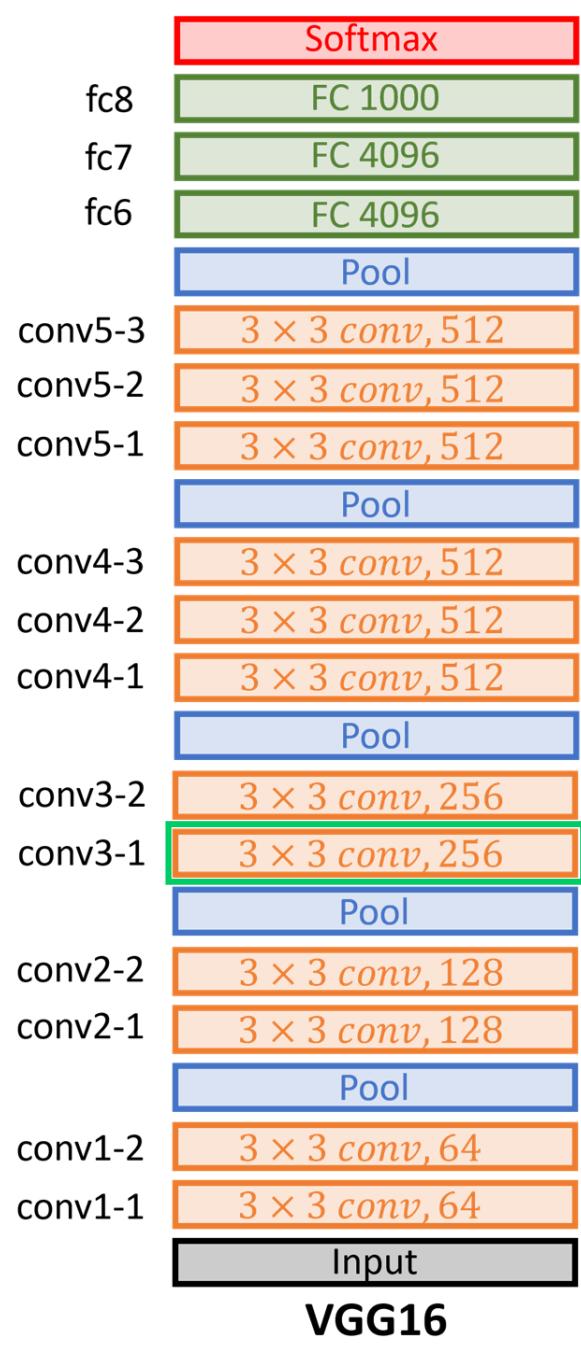


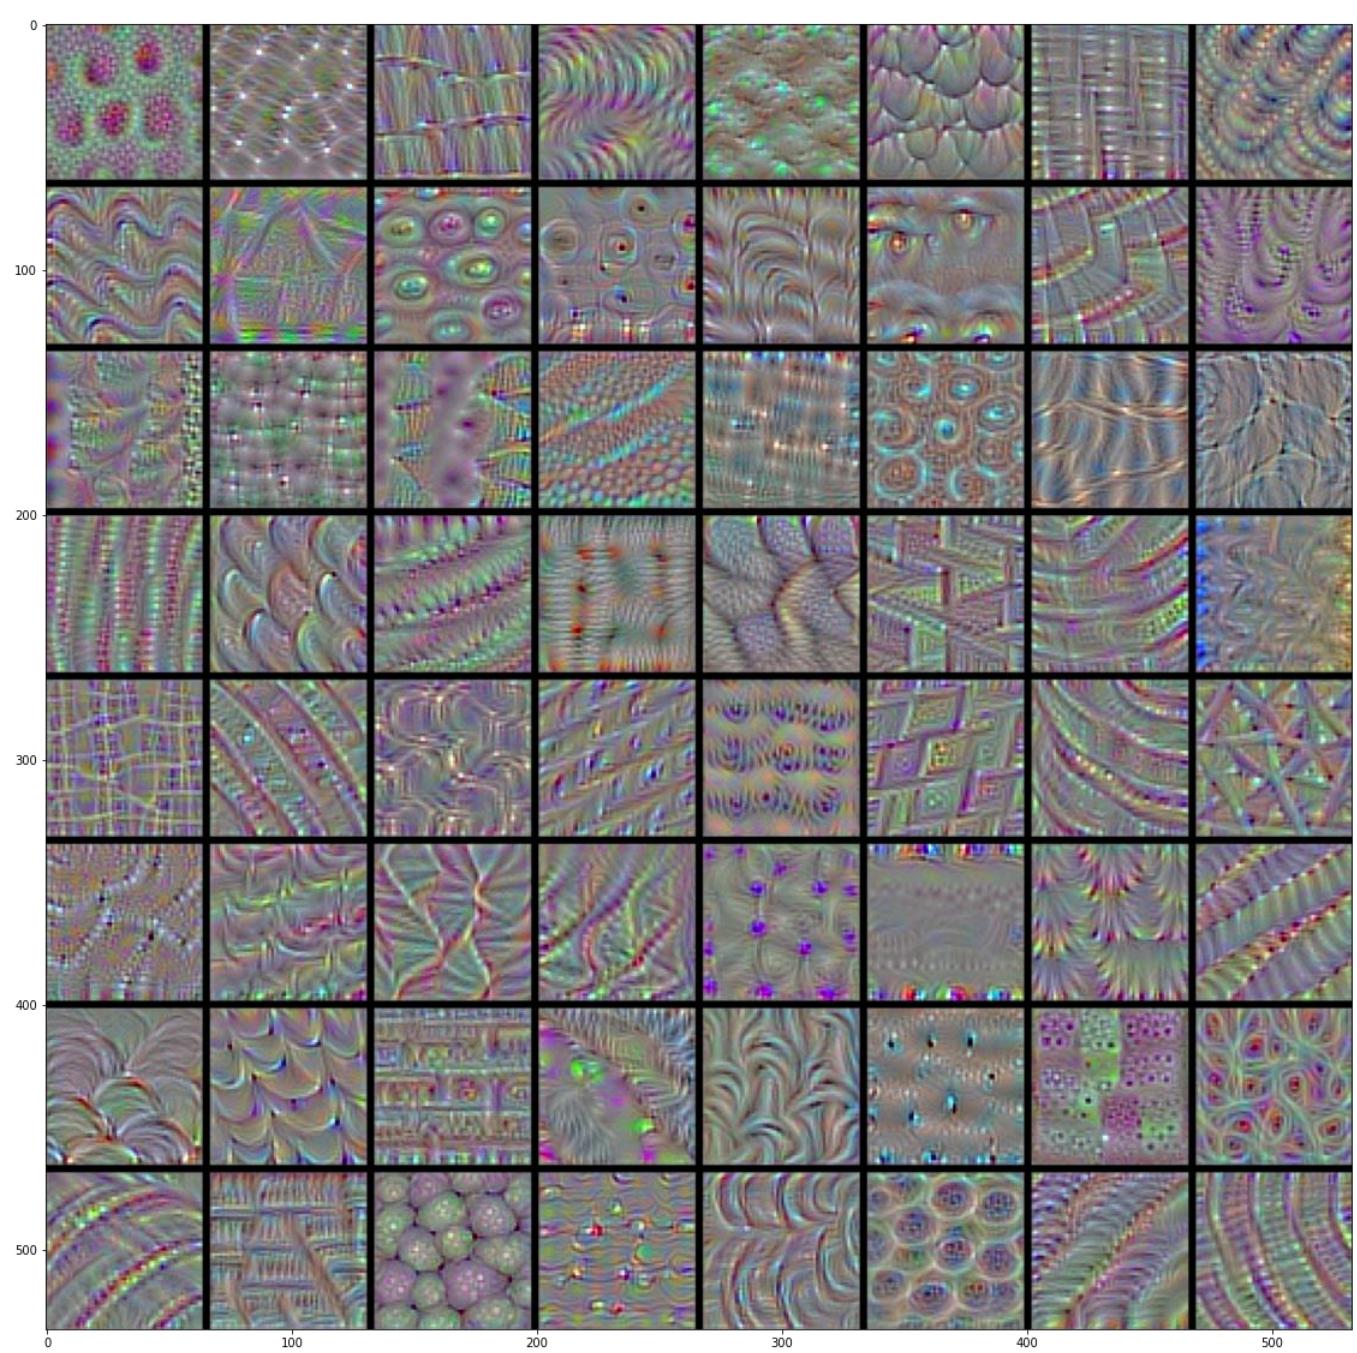
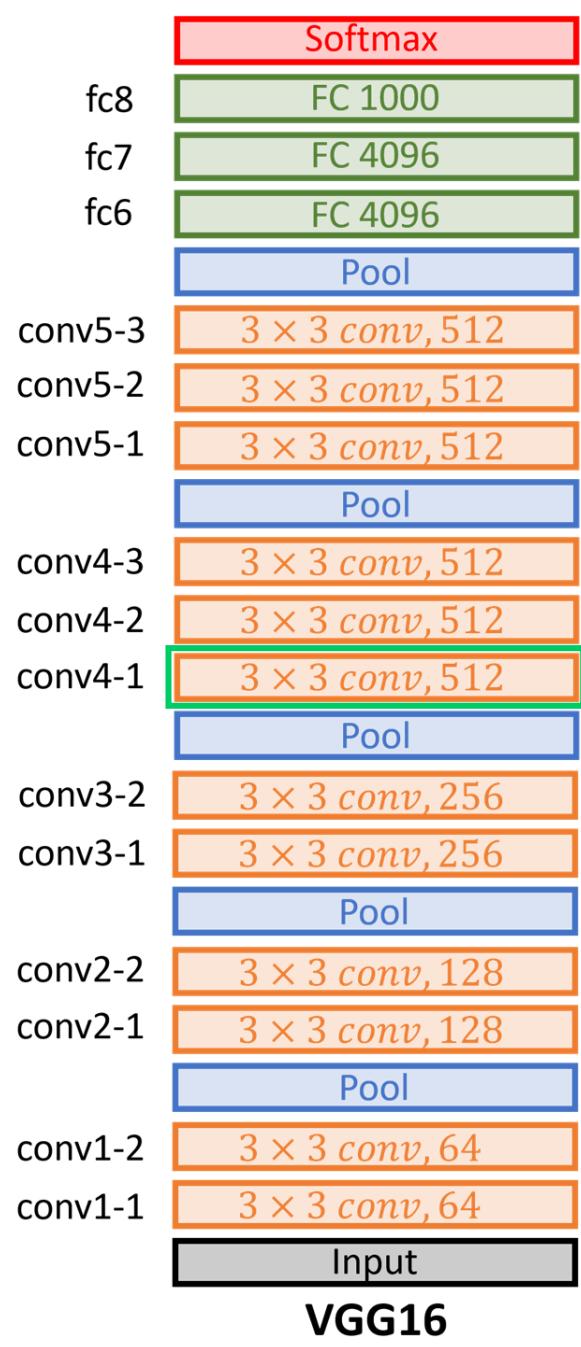
VGG16

64,  $3 \times 3 \times 3$  filters



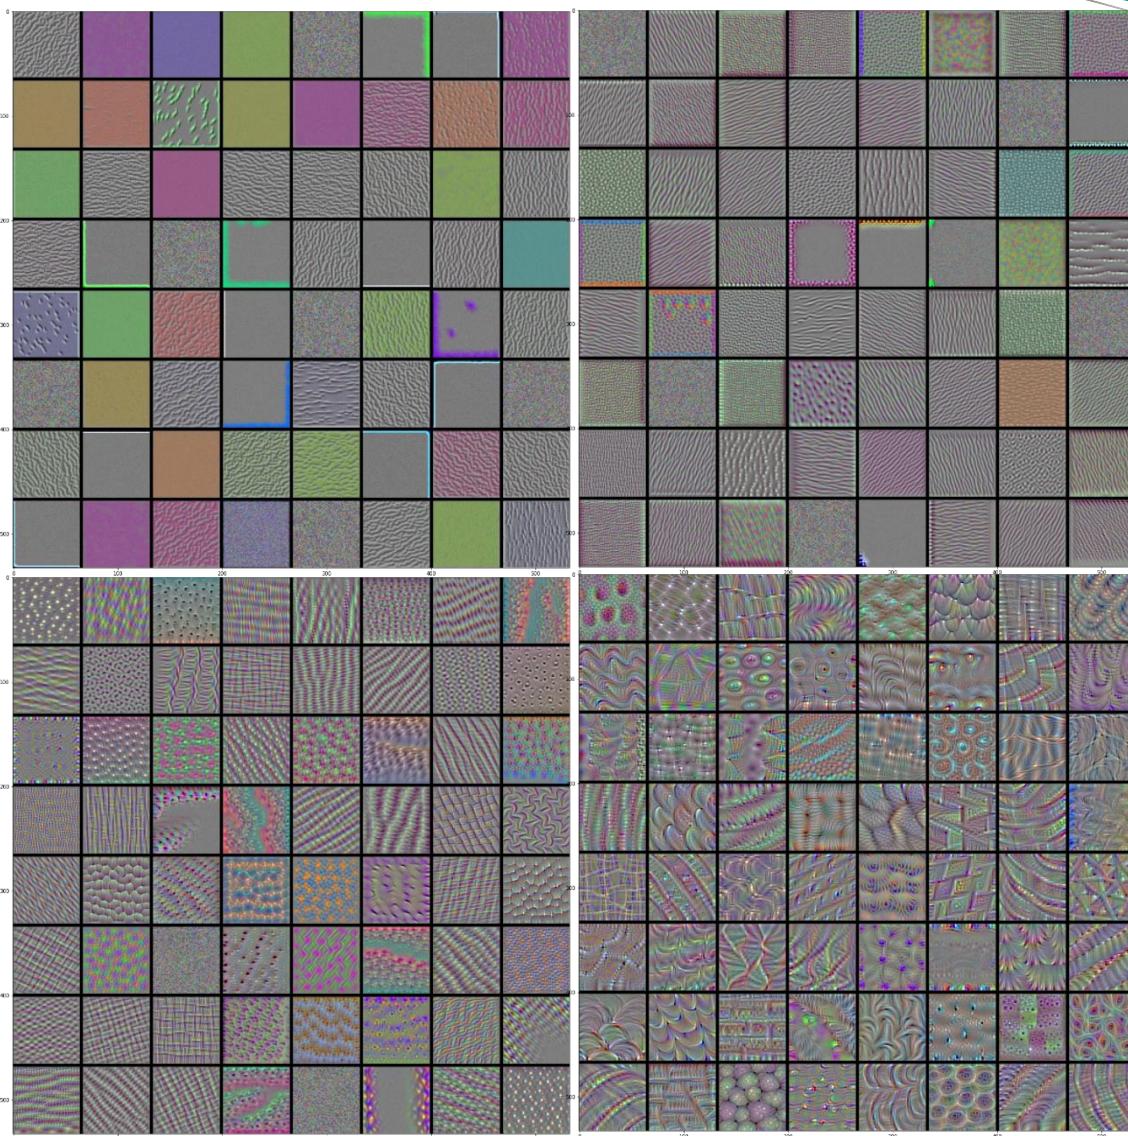






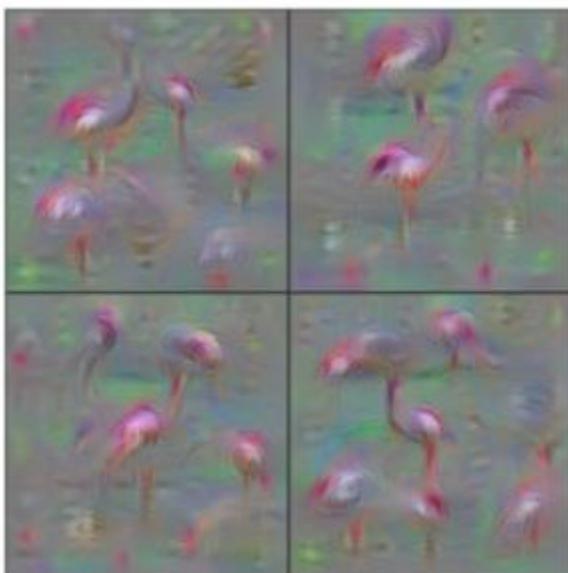
# Visualizing convnet filters

- The filters get increasingly complex and refined as you go higher in the model:
  - The filters from the first layer in the model encode simple directional edges and colors (or colored edges, in some cases)
  - The filters from block2\_conv1 encode simple textures made from combinations of edges and colors
  - The filters in higher layers begin to resemble textures found in natural image: feathers, eyes, leaves, and so on

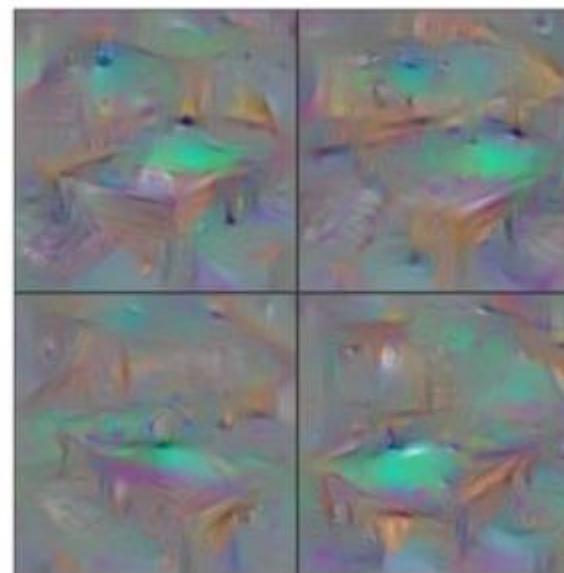


# DeepVis

- Visualizing output neurons
  - The synthetic input image that best fires that neuron



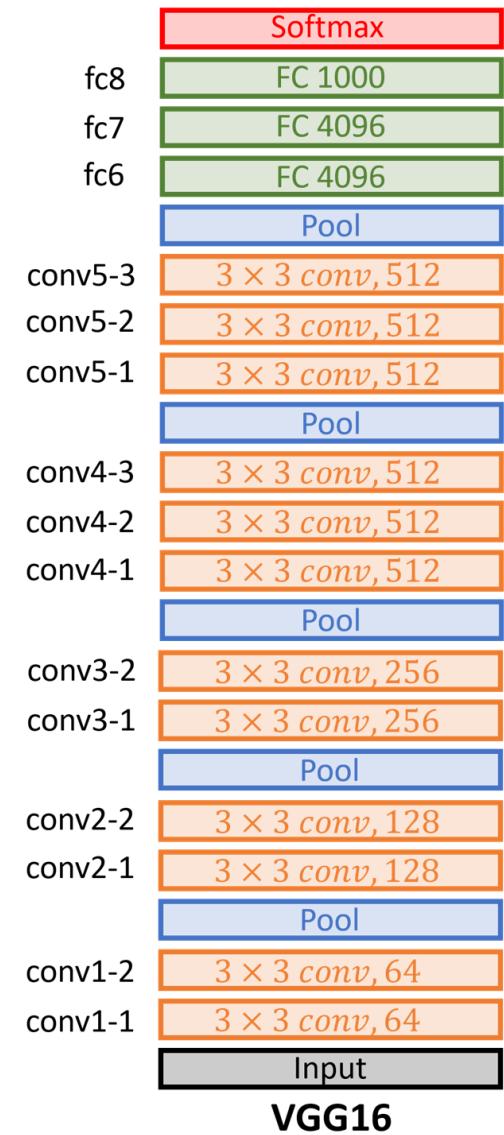
Flamingo



Billiard Table

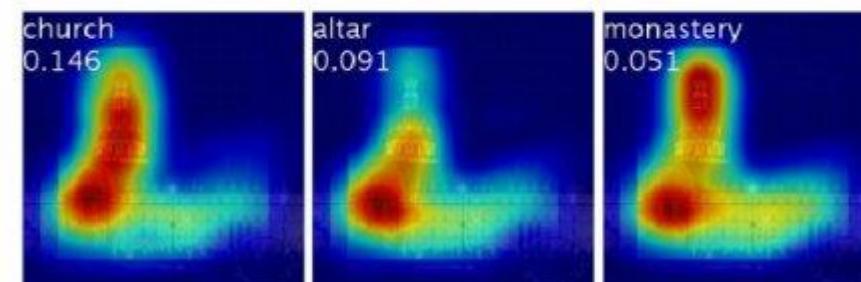
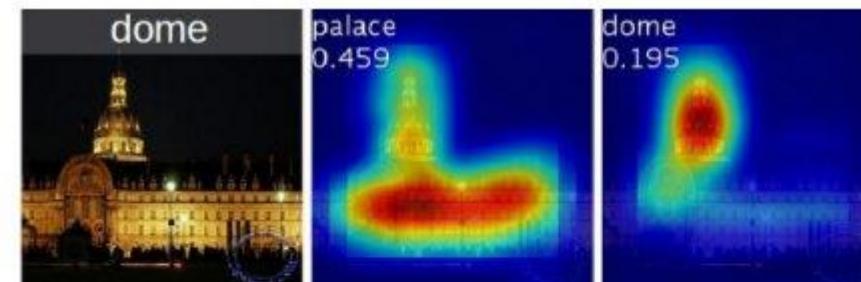


School Bus

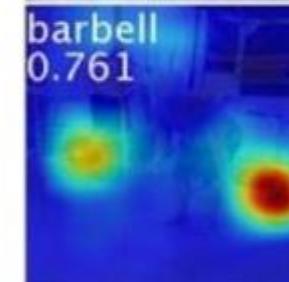


# Visualizing heatmaps of class activation

- Class activation map (CAM) visualization
  - Heatmaps of class activation over input images
- Indicates how important each location is with respect to the class under consideration



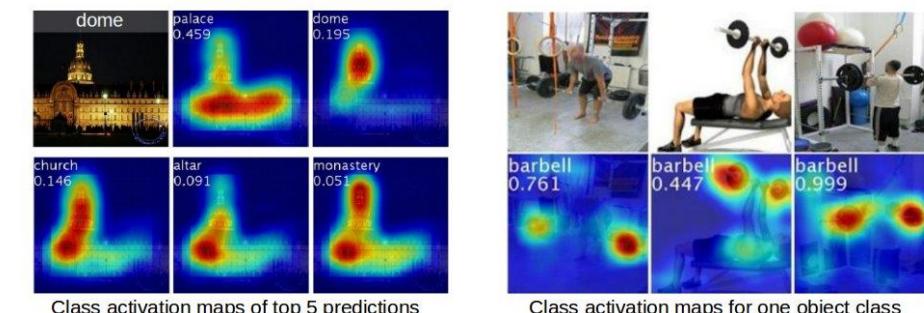
Class activation maps of top 5 predictions



Class activation maps for one object class

# Visualizing heatmaps of class activation

- Class activation map (CAM) visualization
  - Producing heatmaps of class activation over input images
- Useful for understanding which parts of a given image led a convnet to its final classification decision
  - Helpful for debugging the decision process of a convnet, particularly in the case of a classification mistake
- It also allows you to locate specific objects in an image

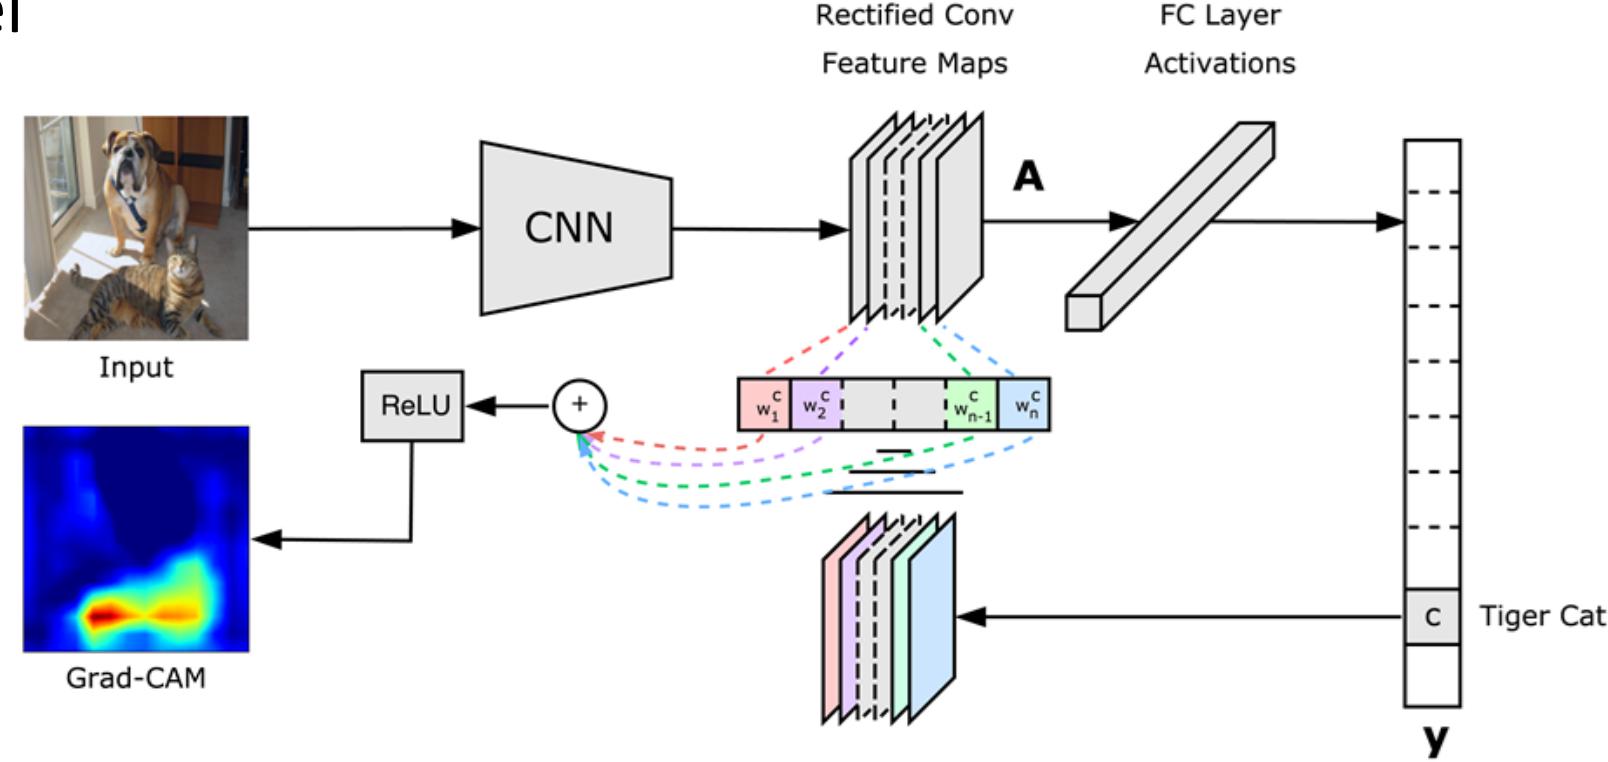


# Grad-CAM

- Take the output feature map of a convolution layer, given an input image, and weight every channel in that feature map by the gradient of the class with respect to the channel

$$w_k^c = \frac{1}{Z} \sum_i \sum_j \frac{\partial y^c}{\partial A_{ij}^k}$$

$$L_{\text{Grad-CAM}}^c = \text{ReLU} \left( \sum_k w_k^c A^k \right)$$



## Listing 5.40 Loading the VGG16 network with pretrained weights

```
from keras.applications.vgg16 import VGG16  
model = VGG16(weights='imagenet')
```

Note that you include the densely connected classifier on top; in all previous cases, you discarded it.

## Listing 5.41 Preprocessing an input image for VGG16

```
from keras.preprocessing import image  
from keras.applications.vgg16 import preprocess_input, decode_predictions  
import numpy as np
```

```
▷ img_path = '/Users/fchollet/Downloads/creative_commons_elephant.jpg'
```

```
▷ img = image.load_img(img_path, target_size=(224, 224))
```

```
x = image.img_to_array(img)
```

float32 Numpy array of shape  
(224, 224, 3)

```
x = np.expand_dims(x, axis=0)
```

Adds a dimension to transform the array  
into a batch of size (1, 224, 224, 3)

```
x = preprocess_input(x)
```

Preprocesses the batch (this does  
channel-wise color normalization)

Python Imaging Library (PIL) image  
of size 224 × 224

Local path to the target image

```
>>> preds = model.predict(x)  
>>> print('Predicted:', decode_predictions(preds, top=3)[0])
```



```
>>> preds = model.predict(x)
>>> print('Predicted:', decode_predictions(preds, top=3) [0])

Predicted: [ (u'n02504458', u'African_elephant', 0.92546833),
(u'n01871265', u'tusker', 0.070257246),
(u'n02504013', u'Indian_elephant', 0.0042589349)]
```

```
>>> np.argmax(preds[0])
```

386



## Listing 5.42 Setting up the Grad-CAM algorithm

“African elephant” entry in the prediction vector

```
► african_elephant_output = model.output[:, 386]
```

Output feature map of the block5\_conv3 layer, the last convolutional layer in VGG16

Gradient of the “African elephant” class with regard to the output feature map of block5\_conv3

```
► grads = K.gradients(african_elephant_output, last_conv_layer.output)[0]
```

Vector of shape (512,), where each entry is the mean intensity of the gradient over a specific feature-map channel

```
pooling_grads = K.mean(grads, axis=(0, 1, 2))
```

```
iterate = K.function([model.input],  
                    [pooling_grads, last_conv_layer.output[0]])
```

```
► pooled_grads_value, conv_layer_output_value = iterate([x])
```

```
for i in range(512):  
    conv_layer_output_value[:, :, i] *= pooled_grads_value[i]
```

```
heatmap = np.mean(conv_layer_output_value, axis=-1)
```

Values of these two quantities, as Numpy arrays, given the sample image of two elephants

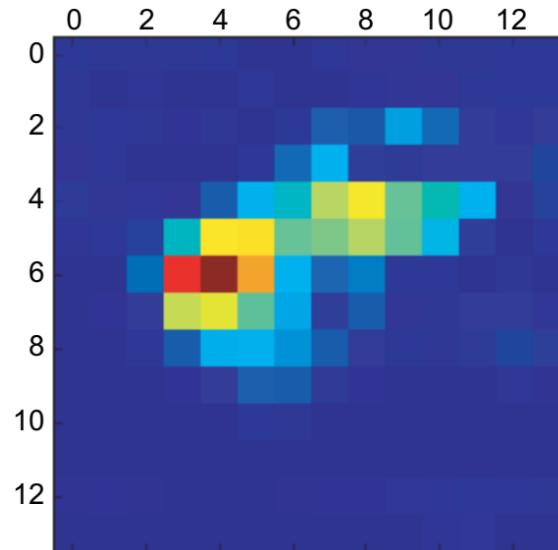
The channel-wise mean of the resulting feature map is the heatmap of the class activation.

Multiplies each channel in the feature-map array by “how important this channel is” with regard to the “elephant” class

Lets you access the values of the quantities you just defined: pooled\_grads and the output feature map of block5\_conv3, given a sample image

## Listing 5.43 Heatmap post-processing

```
heatmap = np.maximum(heatmap, 0)  
heatmap /= np.max(heatmap)  
plt.matshow(heatmap)
```



### Listing 5.44 Superimposing the heatmap with the original picture

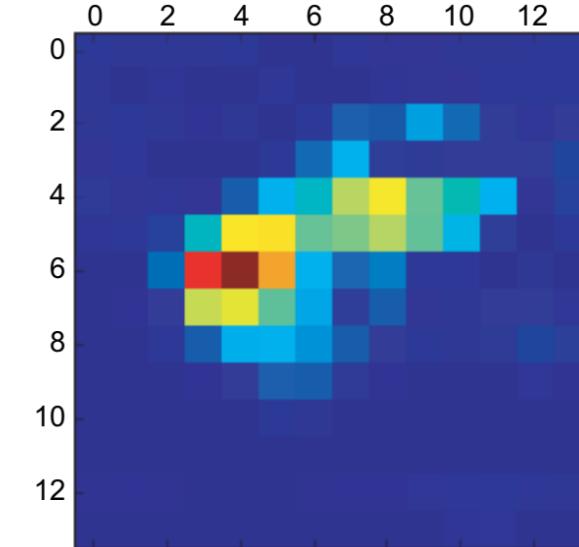
```
import cv2
img = cv2.imread(img_path) ↳ Uses cv2 to load the original image
heatmap = cv2.resize(heatmap, (img.shape[1], img.shape[0])) ↳ Resizes the heatmap to be the same size as the original image
heatmap = np.uint8(255 * heatmap) ↳ Converts the heatmap to RGB
superimposed_img = heatmap * 0.4 + img
cv2.imwrite('/Users/fchollet/Downloads/elephant_cam.jpg', superimposed_img) ↳ 0.4 here is a heatmap intensity factor.
```

Applies the heatmap to the original image



### Listing 5.43 Heatmap post-processing

```
heatmap = np.maximum(heatmap, 0)
heatmap /= np.max(heatmap)
plt.matshow(heatmap)
```



# Visualizing heatmaps of class activation

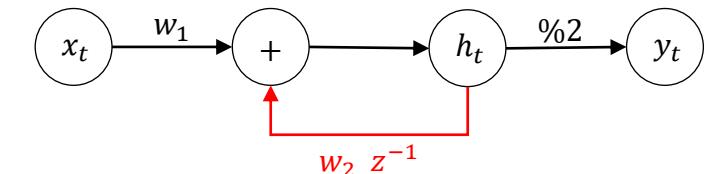
- This visualization technique answers two important questions:
  - Why did the network think this image contained an African elephant?
  - Where is the African elephant located in the picture?
- The ears of the elephant calf are strongly activated:
  - This is probably how the network can tell the difference between African and Indian elephants



# **Sequence Modeling: Recurrent Neural Networks**

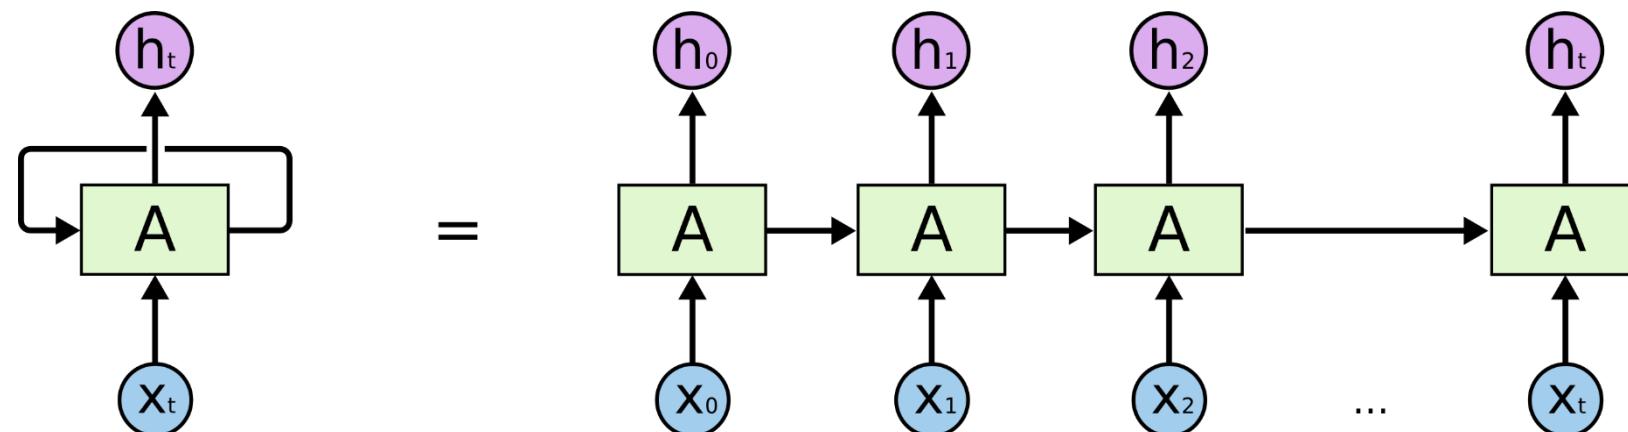
# Motivation

- Not all problems can be converted into one with fixed-length inputs and outputs
- Problems such as Speech Recognition or Time-series Prediction require a system to store and use context information
- Simple case: Output YES if the number of 1s is odd, else NO
  - 1000010101 – No, 100011000000000000 – Yes, ...
- Hard/Impossible to choose a fixed context window



# Recurrent neural networks

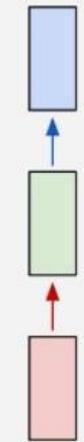
- Recurrent neural networks or RNNs are a family of neural networks for processing sequential data
- RNNs are specialized for processing a sequence of values  $x^{(1)}, \dots, x^{(\tau)}$
- Most recurrent networks can also process sequences of variable length
- An RNN shares the same weights across several time steps
- Computational graphs that include cycles



# Neural Networks

- Dense and ConvNet have no memory!
  - They process the input shown to them independently (with no state in between them)
- These are called feedforward networks

one to one

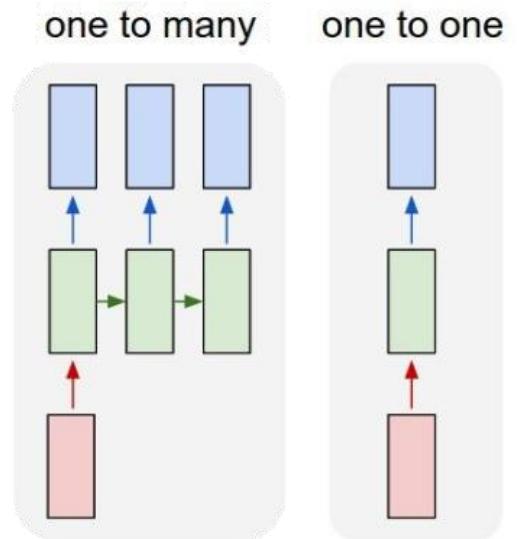


→ Cat

# Neural Networks



*A cat is sitting on a tree branch*



# Neural Networks

## SENTIMENT ANALYSIS



**NEGATIVE**

Totally dissatisfied with the service. Worst customer care ever.



**NEUTRAL**

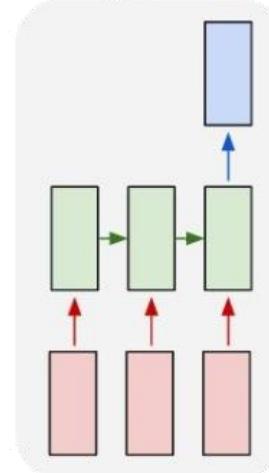
Good Job but I will expect a lot more in future.



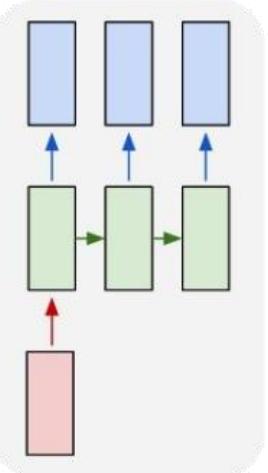
**POSITIVE**

Brilliant effort guys! Loved Your Work.

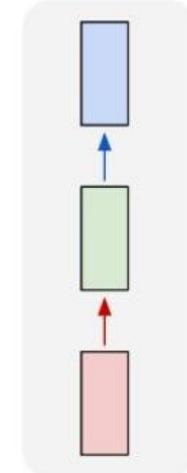
many to one



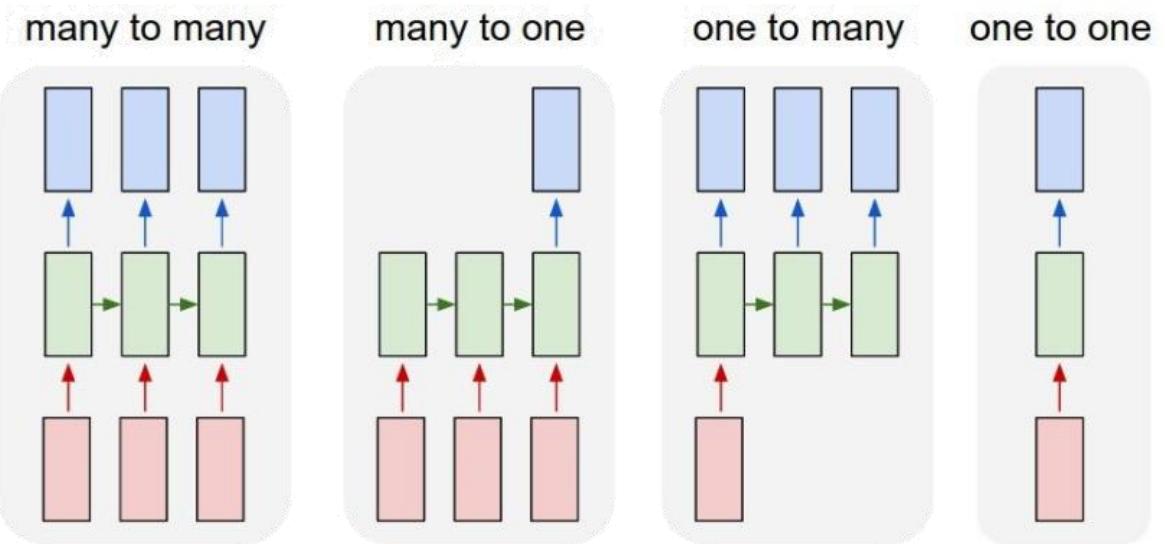
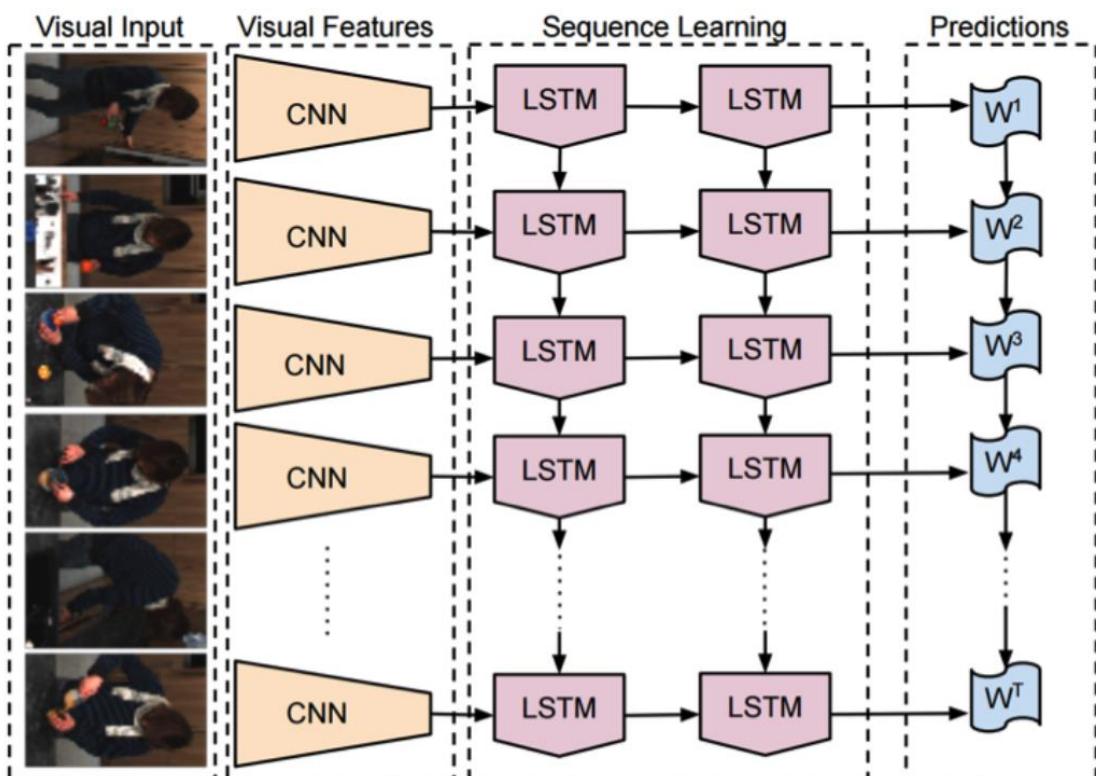
one to many



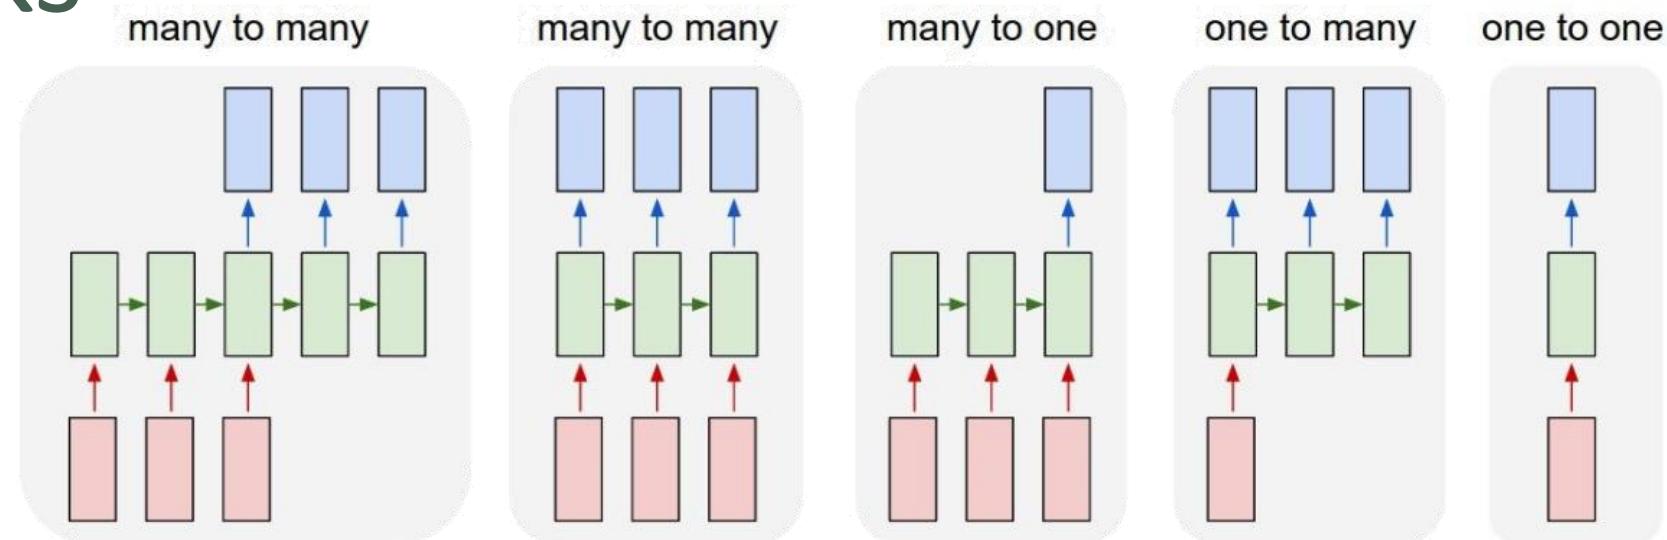
one to one



# Neural Networks



# Neural Networks

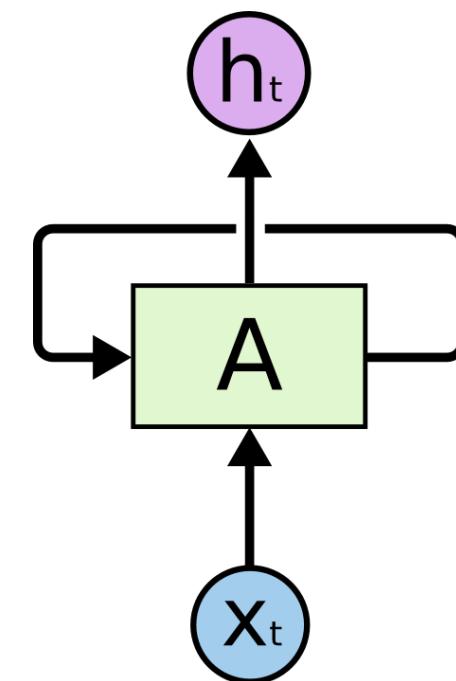


Input: If you face a problem try to find the solution not the reason

Output: مشکل که به وجود اومد بگرد راه حلش را پیدا کن نگردد دنبال این که چرا به وجود اومد

# Recurrent Neural Networks

- RNNs take the previous output or hidden states as inputs
- The composite input at time  $t$  has some historical information about the happenings at time  $T < t$
- RNNs are useful as their intermediate values (state) can store information about past inputs



# Recurrent Neural Networks

$$\mathbf{h}_3 = f_W(\mathbf{h}_2, \mathbf{x}_3)$$

$$= f_W(f_W(\mathbf{h}_1, \mathbf{x}_2), \mathbf{x}_3)$$

$$= f_W(f_W(f_W(\mathbf{h}_0, \mathbf{x}_1), \mathbf{x}_2), \mathbf{x}_3)$$

$$= g^{(3)}(\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3)$$

