

رسالة محمد


Deep Learning

Mohammad Reza Mohammadi

2021

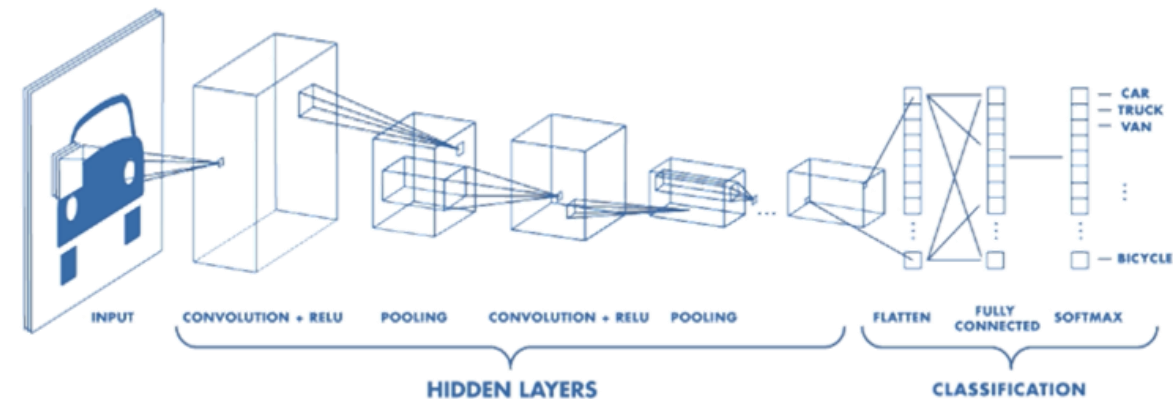
Representation learning

- How learning algorithms share statistical strength across different tasks?
 - Including using information from unsupervised tasks to perform supervised tasks
 - Transfer learned knowledge to tasks for which few or no examples are given but a task representation exists
 - Shared representations are useful to handle multiple modalities or domains




Representation learning

- What makes one representation better than another?
 - A good representation is one that makes a subsequent learning task easier
- We can think of feedforward networks trained by supervised learning as performing a kind of representation learning
 - The last layer of the network is typically a linear classifier
 - The rest of the network learns to provide a representation to this classifier



Representation learning

- Supervised training does not involve explicitly imposing any condition on the learned intermediate features
 - Orthogonal CNNs impose orthogonality on convolutional filters
- We often have very large amounts of unlabeled training data and relatively little labeled training data
 - Semi-Supervised and Self-Supervised Learning
 - Humans and animals are able to learn from very few labeled examples

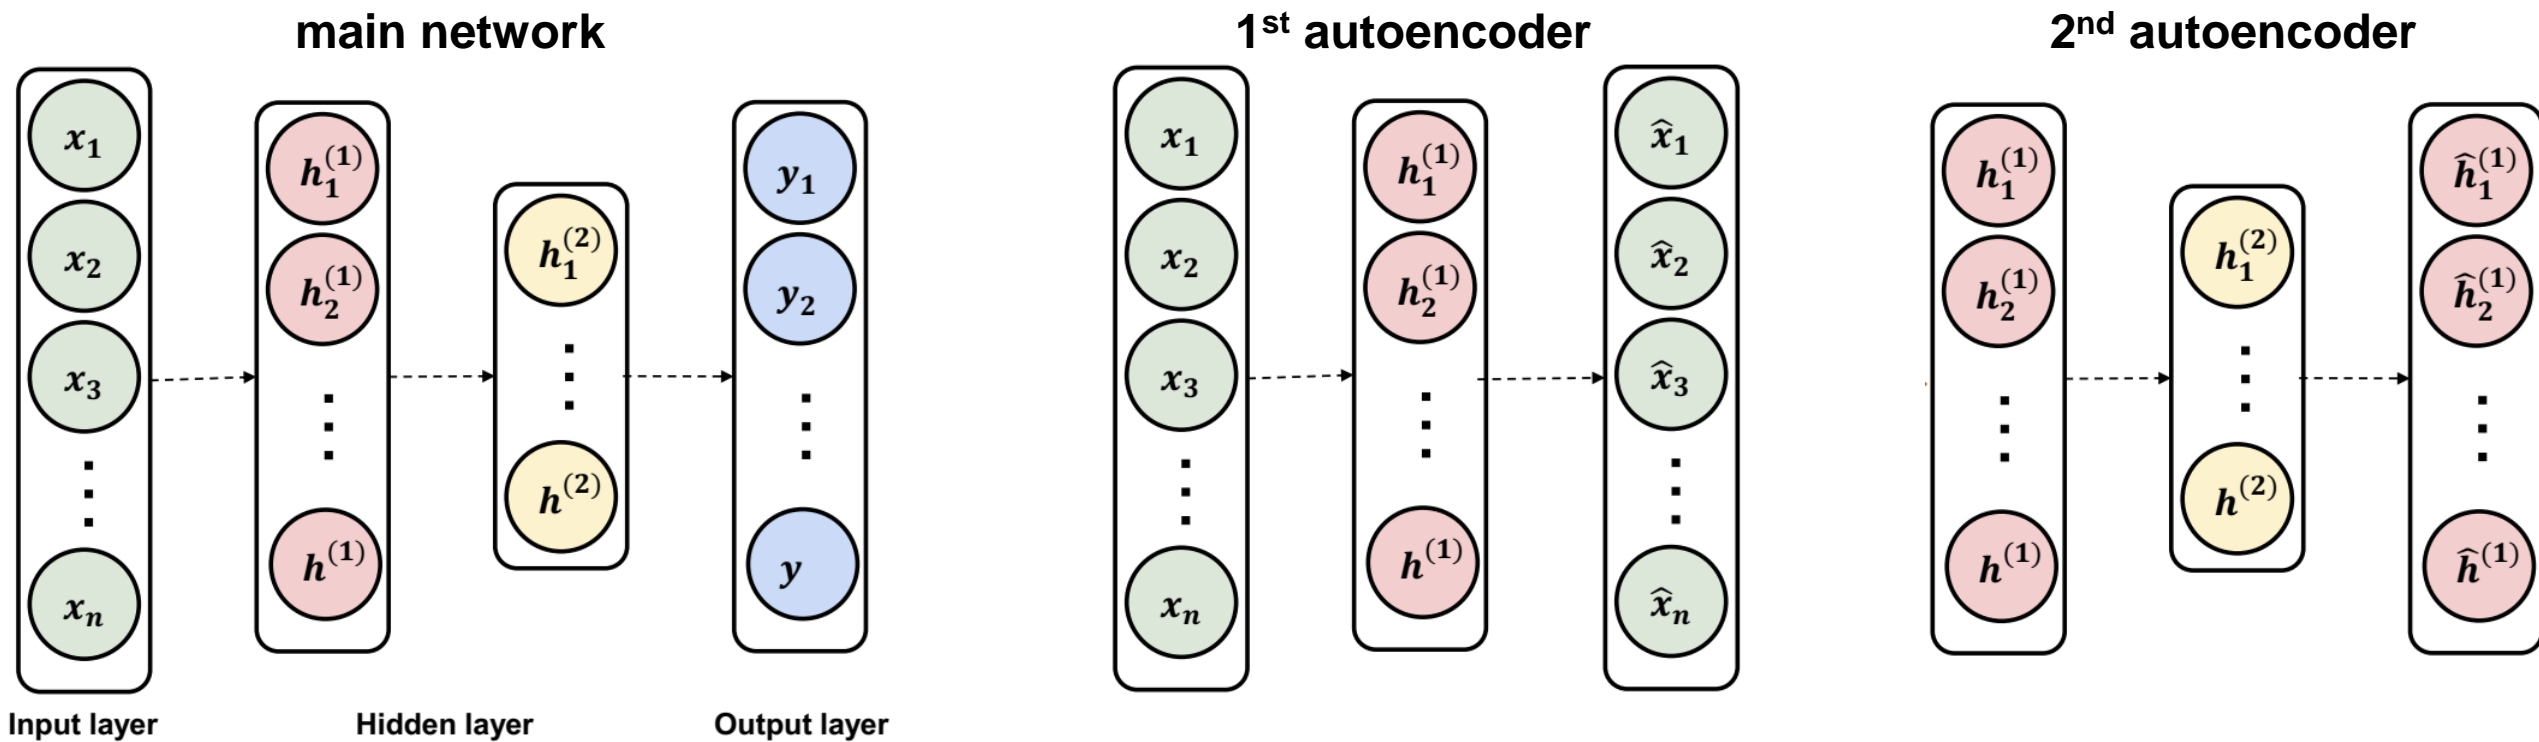


Greedy layer-wise unsupervised pretraining

- Weight initialization is an important design choice when developing deep learning neural network models
- Greedy layer-wise UP have long been used to sidestep the difficulty of jointly training the layers of a deep neural net for a supervised task
- Greedy: UL per a single layer basis, it progresses by layer
- Pretraining: It runs only once before joint training

Greedy layer-wise unsupervised pretraining

- Pretraining proceeds one layer at a time
- It is then possible to perform supervised fine-tuning






Greedy layer-wise unsupervised pretraining

- This approach was performed before the invention and popularization of modern techniques for training very deep networks (ReLU, batch normalization, better optimizers, better architectures, ...)
- Modern approaches typically use simultaneous unsupervised learning and supervised learning rather than two sequential stages
- Most helpful when the number of labeled examples is very small


ResNet

- Such degradation is not caused by overfitting, and adding more layers to a suitably deep model leads to higher training error
- A deeper model should produce no higher training error than its shallower counterpart
- The degradation (of training accuracy) indicates that not all systems are similarly easy to optimize



ResNet

- Let these layers fit a residual mapping $F(x) = H(x) - x$
- We hypothesize that it is easier to optimize the residual mapping than to optimize the original, unreferenced mapping
- To the extreme, if an identity mapping were optimal, it would be easier to push the residual to zero than to fit an identity mapping by a stack of nonlinear layers



Batch normalization (BN)

- Dramatic effect on optimization performance
- Especially for convolutional networks and networks with sigmoidal nonlinearities
- Consider a batch of activations at some layer
 - To make each dimension zero-mean unit-variance, apply:
 - This is a vanilla differentiable function

$$\hat{x}^{(k)} = \frac{x^{(k)} - E[x^{(k)}]}{\sqrt{Var[x^{(k)}]}}$$

Batch normalization

- Compute the empirical mean and variance independently for each dimension

- Input: $x: N \times D$

- Per-channel mean, shape is D :

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{ij}$$

- Per-channel var, shape is D :

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{ij} - \mu_j)^2$$


- Normalized x , shape is $N \times D$

$$\hat{x}_{ij} = \frac{x_{ij} - \mu_j}{\sigma_j}$$

- Output, Shape is $N \times D$

$$y_{ij} = \gamma_j \hat{x}_{ij} + \beta_j$$

Learnable scale and shift



Batch normalization: test time

- Estimates depend on minibatch
 - can't do this at test-time!
- (Running) average of values (μ and σ^2) seen during training
- During testing batchnorm becomes a linear operator!
- Can be fused with the previous fully-connected or conv layer

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{ij}$$

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{ij} - \mu_j)^2$$

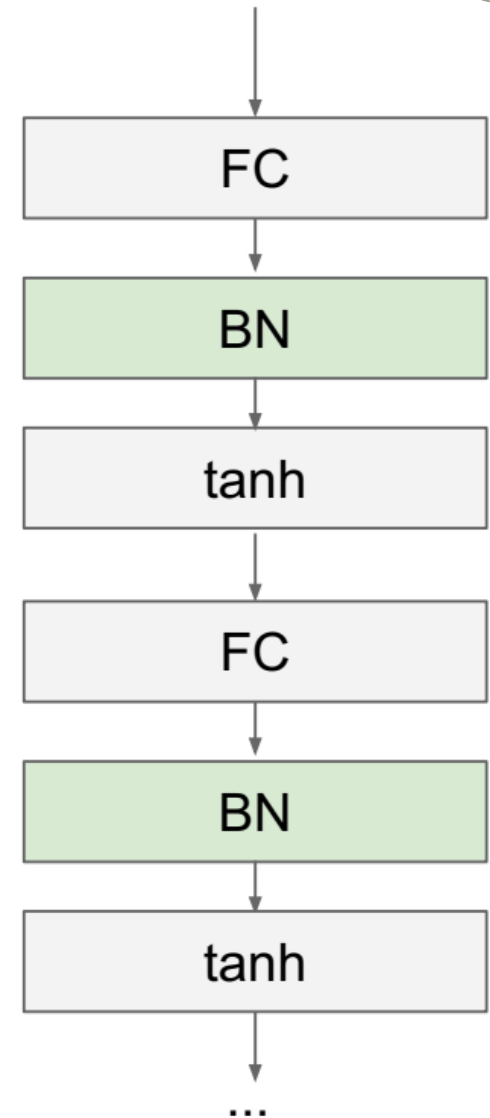
$$\hat{x}_{ij} = \frac{x_{ij} - \mu_j}{\sigma_j}$$

$$y_{ij} = \gamma_j \hat{x}_{ij} + \beta_j$$

Batch normalization

- Usually inserted after Fully Connected or Convolutional layers, and before nonlinearity
- Makes deep networks much easier to train!
- Improves gradient flow
- Allows higher learning rates, faster convergence
- Networks become more robust to initialization
- Acts as regularization during training
- Zero overhead at test-time
 - can be fused with conv!

$$\hat{x}^{(k)} = \frac{x^{(k)} - E[x^{(k)}]}{\sqrt{Var[x^{(k)}]}}$$



Batch normalization for ConvNets

Batch Normalization for
fully-connected layers

$$x: N \times D$$



$$\mu, \sigma: 1 \times D$$

$$\gamma, \beta: 1 \times D$$

$$y = \gamma \frac{x - \mu}{\sigma} + \beta$$

Batch Normalization for
convolutional layers

$$x: N \times W \times H \times C$$



$$\mu, \sigma: 1 \times 1 \times 1 \times C$$

$$\gamma, \beta: 1 \times 1 \times 1 \times C$$

$$y = \gamma \frac{x - \mu}{\sigma} + \beta$$

Layer normalization

Batch Normalization for
fully-connected layers

$$x: N \times D$$



$$\mu, \sigma: 1 \times D$$

$$\gamma, \beta: 1 \times D$$

$$y = \gamma \frac{x - \mu}{\sigma} + \beta$$

Layer Normalization for
fully-connected layers

$$x: N \times D$$



$$\mu, \sigma: N \times 1$$

$$\gamma, \beta: 1 \times D$$

$$y = \gamma \frac{x - \mu}{\sigma} + \beta$$

Instance normalization

Batch Normalization for
convolutional layers

$$x: N \times W \times H \times C$$



$$\mu, \sigma: 1 \times 1 \times 1 \times C$$

$$\gamma, \beta: 1 \times 1 \times 1 \times C$$

$$y = \gamma \frac{x - \mu}{\sigma} + \beta$$

Instance Normalization for
convolutional layers

$$x: N \times W \times H \times C$$




$$\mu, \sigma: N \times 1 \times 1 \times C$$

$$\gamma, \beta: 1 \times 1 \times 1 \times C$$

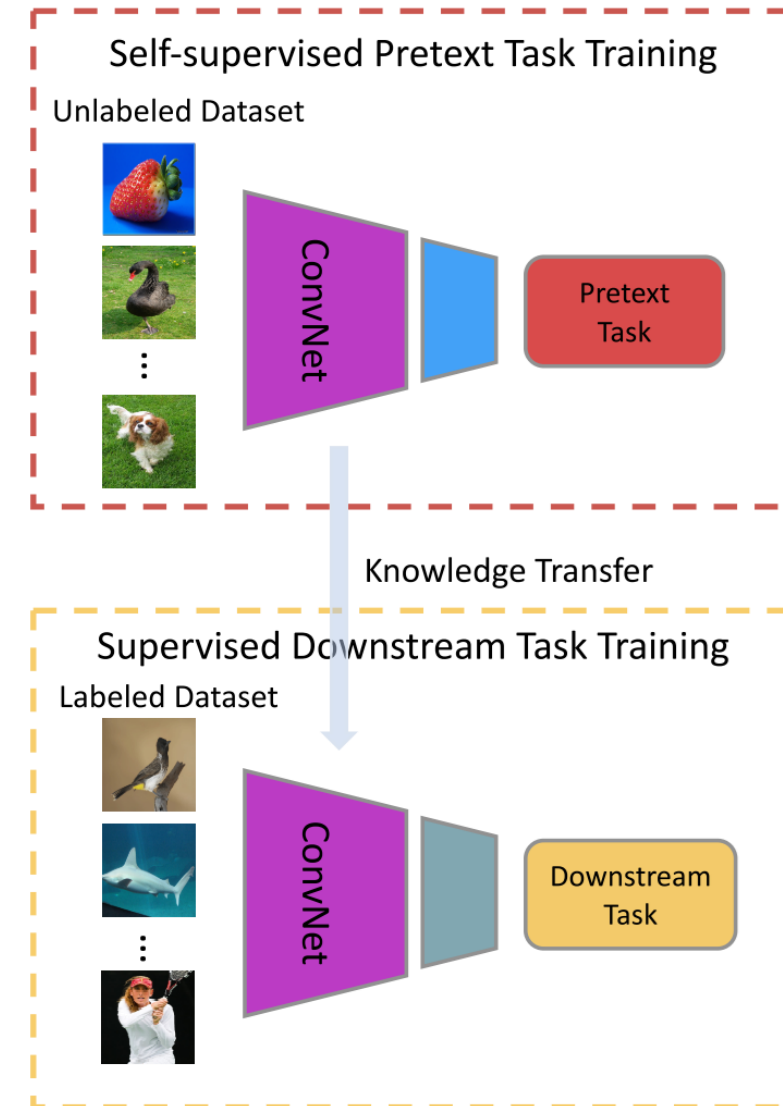
$$y = \gamma \frac{x - \mu}{\sigma} + \beta$$

Comparison of normalization layers




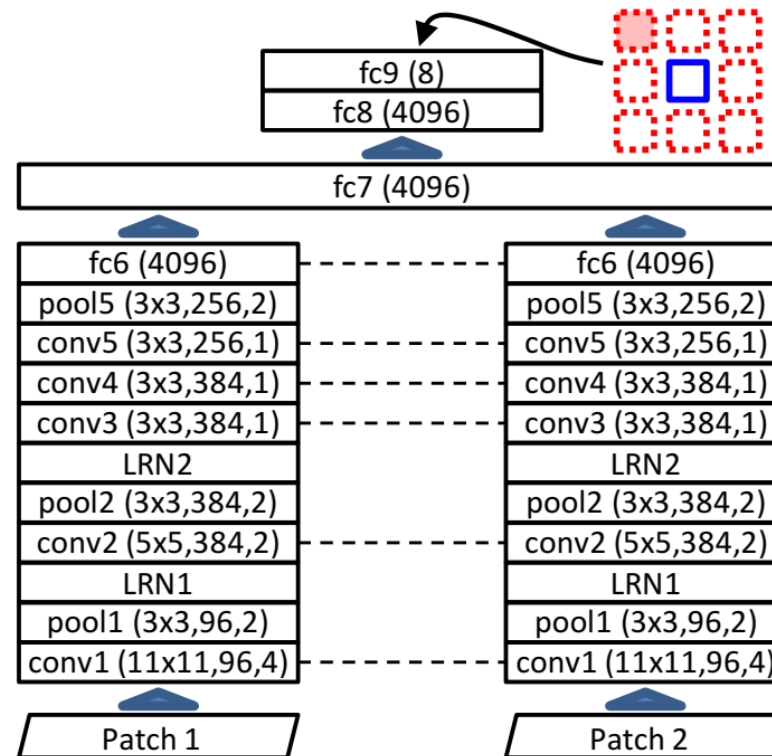
Self-supervised learning

- Self-supervised learning methods are proposed to learn general features from large-scale unlabeled data without using any human-annotated labels
- The pretext tasks share two common properties:
 - Visual features need to be captured by ConvNets to solve the pretext tasks
 - Pseudo labels for the pretext task can be automatically generated based on the attributes of images or videos



Sample pretext: context prediction

- Spatial context as a source of free and plentiful supervisory signal for training a rich visual representation



$$X = \left(\begin{array}{c} \text{cat face} \\ \text{cat ear} \end{array} \right); Y = 3$$

Sample pretext: rotation prediction

- Learn image features by training ConvNets to recognize the 2d rotation

