

رسالة محمد

# Deep Learning

Mohammad Reza Mohammadi  
2021

# Policy

- There are two approaches to train our agent to find this optimal policy  $\pi^*$ :
  - Policy-based methods
    - Directly, teach the agent to learn which action to take, given the state is in
  - Value-based methods
    - Indirectly, teach the agent to learn which state is more valuable and then take the action that leads to the more valuable states

- The link between Value and Policy:

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$




State



$\pi(\text{State}) \rightarrow \text{Action}$

# $\epsilon$ -greedy policy

- In this method agent updates its initial estimates of actions on the basis of received rewards and balances exploration and exploitation by choosing exploratory action with  $\epsilon$  probability and optimal action rest of the time



Initialize, for  $a = 1$  to  $k$ :

$$Q(a) \leftarrow 0$$

$$N(a) \leftarrow 0$$

Repeat forever:

$$A \leftarrow \begin{cases} \arg \max_a Q(a) & \text{with probability } 1 - \epsilon \\ \text{a random action} & \text{with probability } \epsilon \end{cases} \quad (\text{breaking ties randomly})$$


$$R \leftarrow \text{bandit}(A)$$

$$N(A) \leftarrow N(A) + 1$$

$$Q(A) \leftarrow Q(A) + \frac{1}{N(A)} [R - Q(A)]$$

# State-Value function


- The state value function under a policy  $\pi$
- For each state, the state-value function outputs the expected return if the agent starts at that state, and then follow the policy forever after



$$\underbrace{V_\pi(s)}_{\text{Value of state } s} = \underbrace{\mathbf{E}_\pi}_{\text{Expected return}} \left[ \underbrace{G_t}_{\text{And uses the policy to choose its actions for all time steps}} \mid \underbrace{S_t = s}_{\text{If the agent starts at state } s} \right]$$

# Action-Value function

- In the action-value function, for each state and action pair, the action-value function outputs the expected return, if the agent starts in that state and takes the action, and then follows the policy forever after
- The value of taking action  $a$  in state  $s$  under a policy  $\pi$  is:



$$Q_\pi(s, a) = \mathbf{E}_\pi[G_t | S_t = s, A_t = a]$$

Value of state-action pair  $s, a$



Expected return

If the agent starts at state  $s$  and chooses action  $a$

And then uses the policy to choose its actions for all time steps


# Value functions

- In state-value function, we calculate the value of a state ( $S_t$ )
- In action-value function, we calculate the value of state-action pair ( $S_t, A_t$ ) hence the value of taking that action at that state
- Whatever value function we choose (state-value or action-value function), the value is the expected return
- We need to sum all the rewards an agent can get if it starts at that state



# Bellman equation


- The Bellman equation simplifies our value calculation
- With what we learned from now, we know that if we calculate the  $V(S_t)$ , we need to calculate the return starting at that state then follow the policy forever after
- So to calculate  $V(S_t)$  we need to make the sum of the expected rewards





# Bellman equation

- Then, to calculate the  $V(S_{t+1})$ , we need to calculate the return starting at that state  $S_{t+1}$
- That's a quite dull process if you need to do it for each state value or state-action value
- Instead of calculating for each state or each state-action pair the expected return, we can use the Bellman equation



# Bellman equation

- The Bellman equation is a recursive equation
  - Instead of starting for each state from the beginning and calculating the return, we can consider the value of any state as
- The immediate reward ( $R_{t+1}$ ) + the discounted value of the state that follows ( $\gamma * S_{t+1}$ )

$$\underbrace{V_{\pi}(s)}_{\text{Value of state } s} = \underbrace{\mathbf{E}_{\pi}}_{\substack{\text{Expected value of} \\ \text{immediate reward} \\ \text{And uses the policy to} \\ \text{choose its actions for} \\ \text{all time steps}}} \underbrace{[R_{t+1} + \gamma * V_{\pi}(S_{t+1})]}_{\substack{+ \text{ the discounted value of} \\ \text{next\_state}}} \underbrace{|S_t = s]}_{\substack{\text{If the agent} \\ \text{starts at state } s}}$$

# Monte Carlo vs Temporal Difference Learning

- The idea of RL is that using the experience taken, given the reward, it will update its value or its policy
- Monte Carlo and Temporal Difference Learning are two different strategies on how to train our value function or our policy function
  - Both of them use experience to solve the RL problem
- Monte Carlo uses an entire episode of experience before learning
- Temporal Difference uses only a step ( $S_t, A_t, R_{t+1}, S_{t+1}$ ) to learn

# Monte Carlo

- Monte Carlo waits until the end of the episode, then calculates  $G_t$  (return) and uses it as a target for updating  $V(S_t)$ 
  - It requires a complete entire episode of interaction before updating our value function

$$\underline{V(S_t)} \leftarrow \underline{V(S_t)} + \underline{\alpha} [\underline{G_t} - \underline{V(S_t)}]$$

New value of state t

Former estimation of value of state t  
(= Expected return starting at that state)

Learning Rate


Return at timestep t

Former estimation of value of state t  
(= Expected return starting at that state)

# Monte Carlo

$$V(S_t) \leftarrow V(S_t) + \alpha[G_t - V(S_t)]$$


- We always start the episode at the same starting point
- We try actions using our policy
  - For example, epsilon greedy
- We get the reward and the next state
- We terminate if the cat eats us or if we move  $> 10$  steps
  - We have a list of states, actions, rewards, and next states
- The agent will sum the total rewards  $G_t$
- It will then update  $V(S_t)$  based on the formula
- Then start a new game with this new knowledge



# Monte Carlo: Example

$$V(S_t) \leftarrow V(S_t) + \alpha[G_t - V(S_t)]$$

- Consider initial values = 0,  $\alpha = 0.1$ ,  $\gamma = 1$
- $G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots$
- $G_0 = 1 + 0 + 0 + 0 + 0 + 0 + 1 + 1 + 0 + 0 = 3$
- $V(S_0) = 0 + 0.1 [3 - 0] = 0.3$
- $G_1 = 0 + 0 + 0 + 0 + 0 + 1 + 1 + 0 + 0 = 2$
- $V(S_1) = 0 + 0.1 [2 - 0] = 0.2$
  
- $G_9 = 0 = 0$
- $V(S_9) = 0 + 0.1 [0 - 0] = 0$



# Temporal Difference Learning

- The idea is that with TD we update the  $V(S_t)$  at each step
- But because we didn't play during an entire episode, we don't have  $G_t$  (expected return), instead, we estimate  $G_t$  by adding  $R_{t+1}$  and the discounted value of next state


$$V(S_t) \leftarrow V(S_t) + \alpha[G_t - V(S_t)]$$

$$\underbrace{V(S_t)}_{\text{New value of state } t} \leftarrow \underbrace{V(S_t)}_{\text{Former estimation of value of state } t} + \underbrace{\alpha}_{\text{Learning Rate}} [\underbrace{R_{t+1}}_{\text{Reward}} + \underbrace{\gamma V(S_{t+1})}_{\text{Discounted value of next state}} - \underbrace{V(S_t)}_{\text{TD Target}}]$$

# Temporal Difference Learning

- TD waits for only one interaction (one step)  $S_{t+1}$  to form a TD target and update  $V(S_t)$  using  $R_{t+1}$  and  $\gamma V(S_{t+1})$
- We speak about bootstrap because TD bases its update part on an existing estimate  $V(S_{t+1})$  and not a full sample  $G_t$
- This method is called TD(0) or one step TD (update after any individual step)

$$\underbrace{V(S_t)}_{\text{New value of state t}} \leftarrow \underbrace{V(S_t)}_{\text{Former estimation of value of state t}} + \underbrace{\alpha}_{\text{Learning Rate}} [\underbrace{R_{t+1}}_{\text{Reward}} + \underbrace{\gamma V(S_{t+1})}_{\text{Discounted value of next state}} - \underbrace{V(S_t)}_{\text{Former estimation of value of state t}}]$$


  
TD Target



# TD: Example

$$V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

- Consider initial values = 0,  $\alpha = 0.1$ ,  $\gamma = 1$
- It gets a reward  $R_{t+1} = 1$  since it eat a piece of cheese
- $V(S_0) = 0 + 0.1 [1 + 1 \times 0 - 0] = 0.1$
- We just updated our value function for State 0
- Now we continue to interact with this environment with our updated value function



# Monte Carlo vs TD Learning

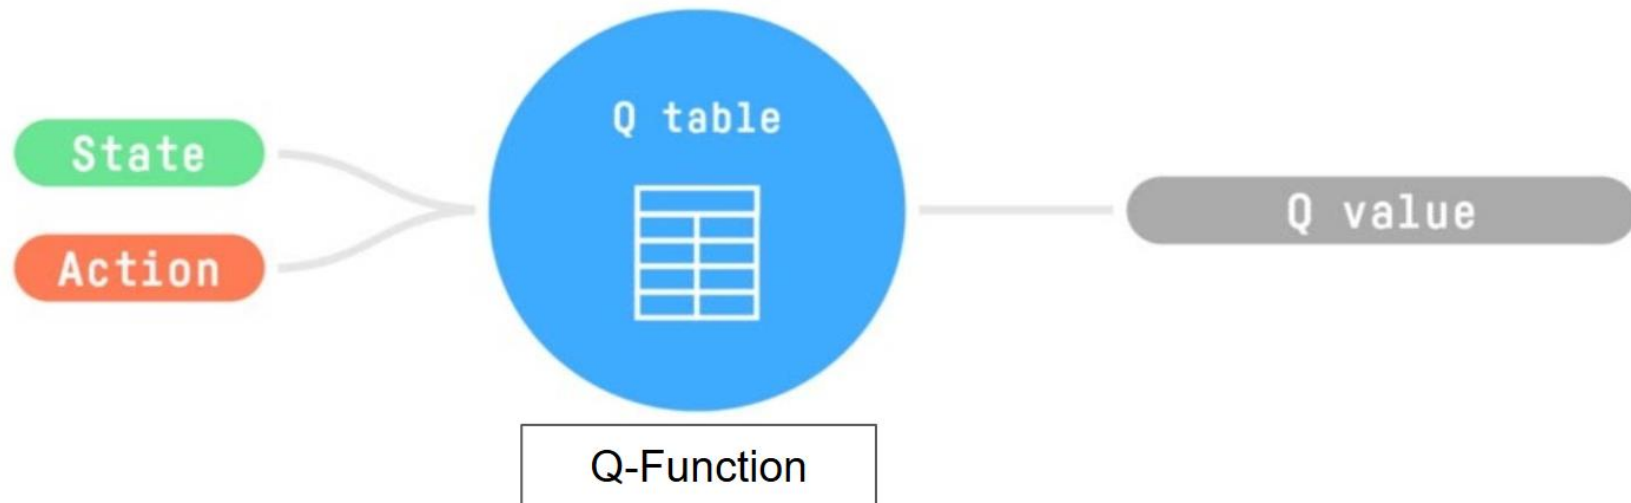
- With Monte Carlo, we update the value function from a complete episode and so we use the actual accurate discounted return of this episode
- With TD learning, we update the value function from a step, so we replace  $G_t$  that we don't have with an estimated return called TD target

$$\text{Monte Carlo: } V(S_t) \leftarrow V(S_t) + \alpha[G_t - V(S_t)]$$

$$\text{TD Learning: } V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

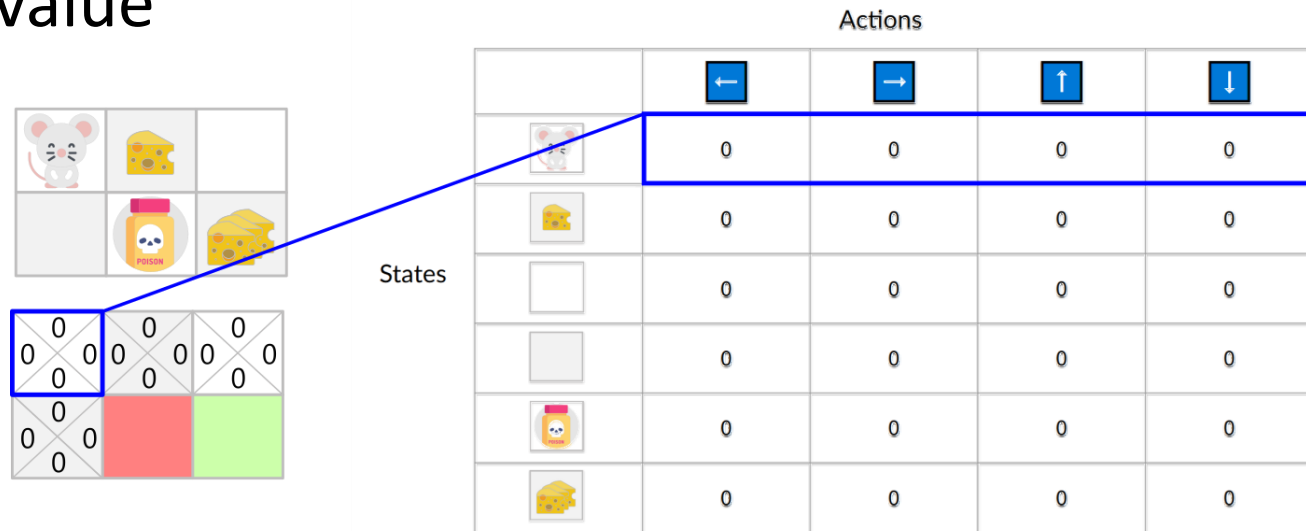
# Q-Learning

- Q-Learning is an off-policy value-based method that uses a TD approach to train its action-value function
- Q-Learning is the algorithm we use to train our Q-Function, an action-value function that determines the value of being at a certain state, and taking a certain action at that state



# Example

- The Q-Table (just initialized that's why all values are = 0), contains for each state, the 4 state-action values
- Q-Function contains a Q-table that contains the value of each state-action
- Given a state and action, our Q-Function will search inside its Q-table to output the value













		Actions			
States		0	0	0	0
		0	0	0	0
		0	0	0	0
		0	0	0	0
		0	0	0	0
		0	0	0	0













# Example

- In the beginning, our Q-Table is useless since it gives arbitrary value for each state-action pair (most of the time we initialize the Q-Table to 0 values)
- But, as we'll explore the environment and update our Q-Table it will give us better and better approximations

				
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0



				
	0	10.8	0	0
	0	9.9	0	-10
	0	0	0	10
	-10	0	0	0
	0	0	0	0
	0	0	0	0



# Q-Learning pseudocode

---

**Algorithm 14:** Sarsamax (Q-Learning)

---

**Input:** policy  $\pi$ , positive integer  $num\_episodes$ , small positive fraction  $\alpha$ , GLIE  $\{\epsilon_i\}$

**Output:** value function  $Q$  ( $\approx q_\pi$  if  $num\_episodes$  is large enough)

Initialize  $Q$  arbitrarily (e.g.,  $Q(s, a) = 0$  for all  $s \in \mathcal{S}$  and  $a \in \mathcal{A}(s)$ , and  $Q(terminal-state, \cdot) = 0$ )

**for**  $i \leftarrow 1$  **to**  $num\_episodes$  **do**

$\epsilon \leftarrow \epsilon_i$

    Observe  $S_0$

$t \leftarrow 0$

**repeat**

        Choose action  $A_t$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)

        Take action  $A_t$  and observe  $R_{t+1}, S_{t+1}$

$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t))$

$t \leftarrow t + 1$

**until**  $S_t$  is terminal;

**end**

**return**  $Q$

---

# Q-Learning

---

## Algorithm 14: Sarsamax (Q-Learning)

---

**Input:** policy  $\pi$ , positive integer  $num\_episodes$ , small positive fraction  $\alpha$ , GLIE  $\{\epsilon_i\}$

**Output:** value function  $Q$  ( $\approx q_\pi$  if  $num\_episodes$  is large enough)

Initialize  $Q$  arbitrarily (e.g.,  $Q(s, a) = 0$  for all  $s \in \mathcal{S}$  and  $a \in \mathcal{A}(s)$ , and  $Q(\text{terminal-state}, \cdot) = 0$ )

**for**  $i \leftarrow 1$  **to**  $num\_episodes$  **do**

$\epsilon \leftarrow \epsilon_i$

    Observe  $S_0$

$t \leftarrow 0$

**repeat**

        Choose action  $A_t$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)

        Take action  $A_t$  and observe  $R_{t+1}, S_{t+1}$

$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t))$











$t \leftarrow t + 1$

**until**  $S_t$  is terminal;

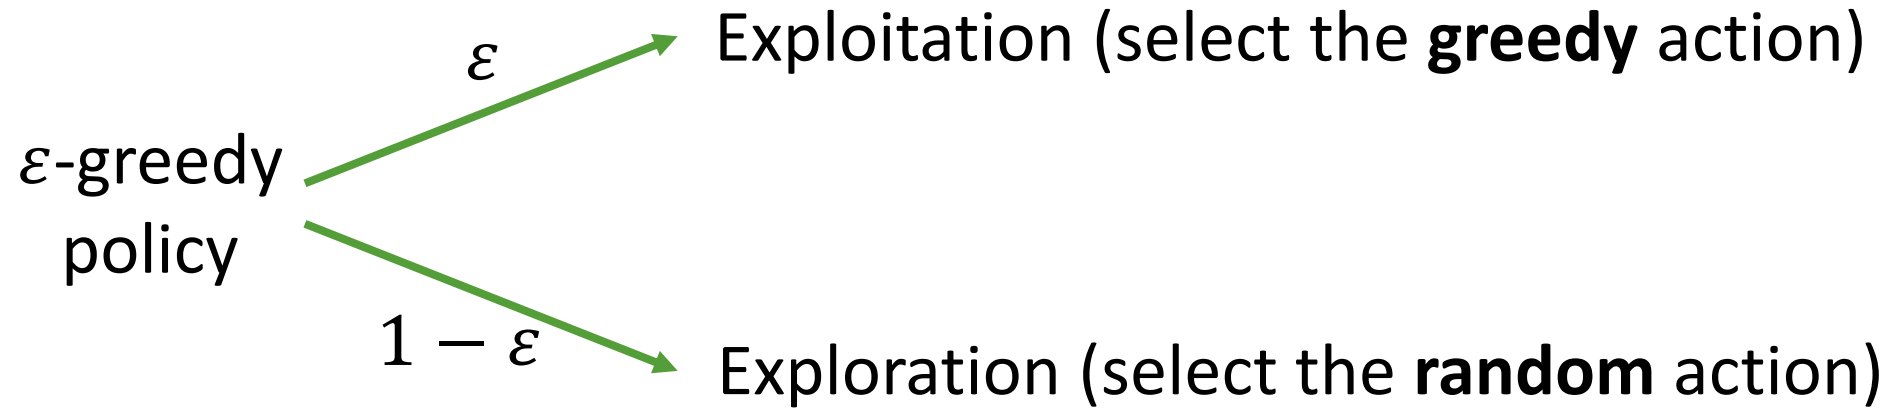
**end**

**return**  $Q$

---

				
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0

# Q-Learning



---

**Algorithm 14:** Sarsamax (Q-Learning)

---

**Input:** policy  $\pi$ , positive integer  $num\_episodes$ , small positive fraction  $\alpha$ , GLIE  $\{\epsilon_i\}$

**Output:** value function  $Q$  ( $\approx q_\pi$  if  $num\_episodes$  is large enough)

Initialize  $Q$  arbitrarily (e.g.,  $Q(s, a) = 0$  for all  $s \in \mathcal{S}$  and  $a \in \mathcal{A}(s)$ , and  $Q(\text{terminal-state}, \cdot) = 0$ )

**for**  $i \leftarrow 1$  **to**  $num\_episodes$  **do**

$\epsilon \leftarrow \epsilon_i$

    Observe  $S_0$

$t \leftarrow 0$

**repeat**

        Choose action  $A_t$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)

        Take action  $A_t$  and observe  $R_{t+1}, S_{t+1}$

$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t))$


$t \leftarrow t + 1$

**until**  $S_t$  is terminal;

**end**

**return**  $Q$

---





# Q-Learning

- To update  $Q(S_t, A_t)$ , we need  $S_t, A_t, R_{t+1}, S_{t+1}$
- We use  $R_{t+1}$  and to get the best next-state-action pair value, we select with a greedy-policy (so not our epsilon greedy policy) the next best action

---

## Algorithm 14: Sarsamax (Q-Learning)

---

**Input:** policy  $\pi$ , positive integer  $num\_episodes$ , small positive fraction  $\alpha$ , GLIE  $\{\epsilon_i\}$

**Output:** value function  $Q$  ( $\approx q_\pi$  if  $num\_episodes$  is large enough)

Initialize  $Q$  arbitrarily (e.g.,  $Q(s, a) = 0$  for all  $s \in \mathcal{S}$  and  $a \in \mathcal{A}(s)$ , and  $Q(\text{terminal-state}, \cdot) = 0$ )

**for**  $i \leftarrow 1$  **to**  $num\_episodes$  **do**

$\epsilon \leftarrow \epsilon_i$

    Observe  $S_0$

$t \leftarrow 0$

**repeat**

        Choose action  $A_t$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)

        Take action  $A_t$  and observe  $R_{t+1}, S_{t+1}$

$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t))$

$t \leftarrow t + 1$

**until**  $S_t$  is terminal;

**end**

**return**  $Q$

---

$$V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

New value  
of state  $t$

Former  
estimation of  
value of state  
 $t$

Learning  
Rate

Reward

Discounted value of next  
state

TD Target

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

New  
Q-value  
estimation

Former  
Q-value  
estimation

Learning  
Rate

Immediate  
Reward

Discounted Estimate  
optimal Q-value  
of next state

Former  
Q-value  
estimation

TD Target

TD Error

# Off-policy vs On-policy

- Off-policy: using a different policy for **acting** and **updating**
- On-policy: using the same policy for **acting** and **updating**

---

**Algorithm 14:** Sarsamax (Q-Learning)

---

**Input:** policy  $\pi$ , positive integer  $num\_episodes$ , small positive fraction  $\alpha$ , GLIE  $\{\epsilon_i\}$   
**Output:** value function  $Q$  ( $\approx q_\pi$  if  $num\_episodes$  is large enough)  
Initialize  $Q$  arbitrarily (e.g.,  $Q(s, a) = 0$  for all  $s \in \mathcal{S}$  and  $a \in \mathcal{A}(s)$ , and  $Q(\text{terminal-state}, \cdot) = 0$ )  
**for**  $i \leftarrow 1$  **to**  $num\_episodes$  **do**  
     $\epsilon \leftarrow \epsilon_i$   
    Observe  $S_0$   
     $t \leftarrow 0$   
    **repeat**  
        Choose action  $A_t$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)  
        Take action  $A_t$  and observe  $R_{t+1}, S_{t+1}$   
         $Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t))$   
         $t \leftarrow t + 1$   
    **until**  $S_t$  is terminal;  
**end**  
**return**  $Q$

---

---

**Algorithm 13:** Sarsa

---

**Input:** policy  $\pi$ , positive integer  $num\_episodes$ , small positive fraction  $\alpha$ , GLIE  $\{\epsilon_i\}$   
**Output:** value function  $Q$  ( $\approx q_\pi$  if  $num\_episodes$  is large enough)  
Initialize  $Q$  arbitrarily (e.g.,  $Q(s, a) = 0$  for all  $s \in \mathcal{S}$  and  $a \in \mathcal{A}(s)$ , and  $Q(\text{terminal-state}, \cdot) = 0$ )  
**for**  $i \leftarrow 1$  **to**  $num\_episodes$  **do**  
     $\epsilon \leftarrow \epsilon_i$   
    Observe  $S_0$   
    Choose action  $A_0$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)  
     $t \leftarrow 0$   
    **repeat**  
        Take action  $A_t$  and observe  $R_{t+1}, S_{t+1}$   
        Choose action  $A_{t+1}$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)  
         $Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$   
         $t \leftarrow t + 1$   
    **until**  $S_t$  is terminal;  
**end**  
**return**  $Q$

---